

NewSQL

Chapter-2: NoSQL Data Management

Mr. Janmejay Kumar Vishwakarma
Assistant Professor
Department of CSE

Content

1. Introduction to NoSQL.....	1
2. Why NoSQL?.....	3
3. CAP Theorem.....	4
4. BASE Model.....	8
5. Types of NoSQL.....	10
6. Advantages of NoSQL.....	20

NewSQL

- NewSQL is a class of modern relational database management.
- Support SQL and ACID transactions for online transaction processing (**OLTP**) read-write workloads.
 - Have **many small, fast transactions**
 - Access only a small part** of the data using indexes
 - Avoid full table scans** or big joins
 - Run **repetitive queries** with different input values

Key Features/ Advantages

SQL support (like traditional RDBMS).

ACID compliance (ensures data reliability).

High scalability (like NoSQL).

Optimized for **read-write intensive operations**

NewSQL

Why NewSQL?

Traditional RDBMS: Reliable but not scalable.

NoSQL: Scalable but compromises on consistency.

NewSQL offers the best of both worlds.

Popular NewSQL Databases:

- Google Spanner(Paypal, Spotify, Snapchat, X)
- CockroachDB(Netflix, Cisco)
- VoltDB (Huawei, Dream11)
- TiDB (Shopee, Bank of Beijing)

Use Cases

- E-commerce platforms
- Banking and finance systems
- Real-time analytics
- Large-scale web applications

Differences: SQL vs NoSQL vs NewSQL

Feature	SQL (Relational)	NoSQL (Non-relational)	NewSQL
Data Model	Relational tables (fixed schema)	Flexible: Key-value, document, column, graph (flexible)	Relational (adapted for distributed systems)
Scalability	Vertical (scale-up)	Horizontal (scale-out)	Horizontal + Vertical
ACID Support	Full ACID	Limited (BASE, eventual consistency)	Full ACID in distributed setup
SQL Support	Full standard SQL	No standard SQL	Full SQL + distributed extensions
Use Case	OLTP (small-scale, reliable)	Big Data, real-time, flexible workloads	Distributed OLTP with consistency & speed

Map-Reduce

- **MapReduce** is a programming model and an associated implementation for processing and generating big data sets with a parallel and distributed algorithm on a cluster.
- It is **highly scalable and reliable** and is **based on the principle of divide-and-conquer**, which provides built-in fault tolerance and redundancy.
- The MapReduce programming model was developed at Google in 2004 by Jeffrey Dean and Sanjay Ghemawat.

Map-Reduce Tasks

- It involves two key stages: **mapping** and **reducing**.

Mapping: function that processes input data and emits intermediate key-value pairs.

Reducing: function that aggregates or summarizes all values associated with the same key.

- MapReduce operations along with their individual stages:

Map task

- map
- combine (optional)
- partition

Reduce task

- shuffle and sort
- reduce

Map-Reduce Tasks

Map Tasks (Executed on Mapper Nodes)

1.Map: The main transformation stage.

- Input data is split into records and processed to generate intermediate **key-value pairs**.

Example: "apple orange apple" → ("apple", 1), ("orange", 1), ("apple", 1)

2.Combine (Optional):

- A **mini-reducer** that runs **locally** on the mapper node.
- Performs partial aggregation to reduce the amount of data transferred across the network.

Example: Combine ("apple", 1), ("apple", 1) → ("apple", 2) on the same node.

3.Partition:

- Decides **which reducer** each key-value pair should go to.
- Ensures all values for a given key are sent to the same reducer.
- Often uses a **hash function** on the key.

Map-Reduce Tasks

Reduce Tasks (Executed on Reducer Nodes)

1.Shuffle and Sort:

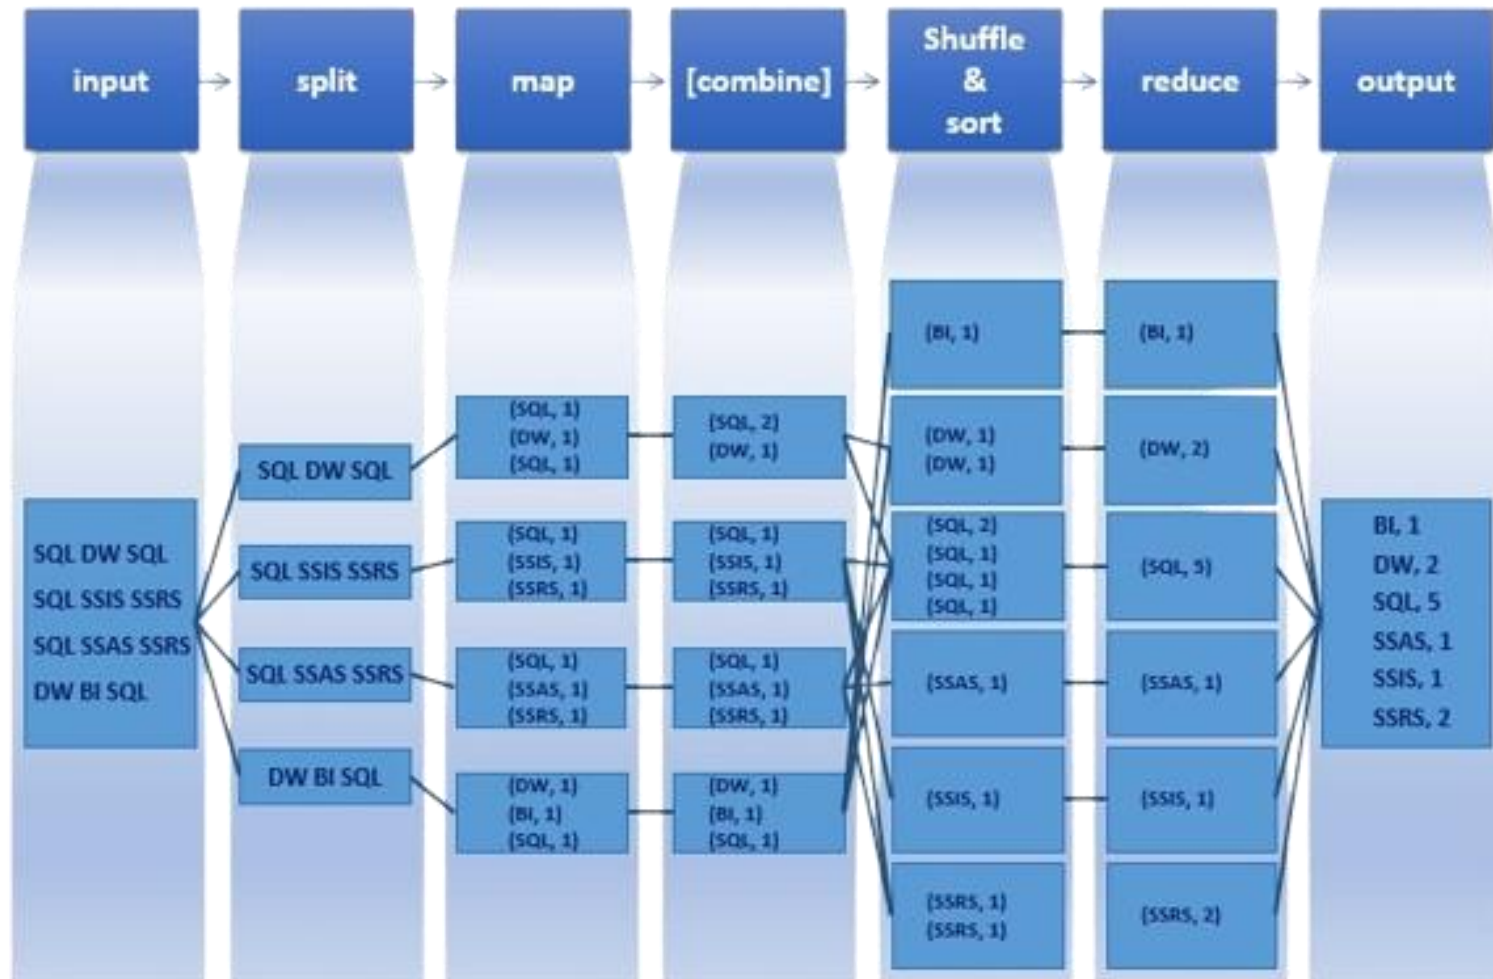
- Intermediate key-value pairs from all mappers are **shuffled** (transferred) across the network to appropriate reducers.
- The keys are then **sorted** before being passed to the reduce function.
- Ensures that a reducer gets all values for a single key grouped together.

2.Reduce:

- Applies user-defined logic on each group of values for a key.
- Produces the **final key-value output**.
- Example: ("apple", [1, 1, 2]) → ("apple", 4)

Example

MapReduce – Word Count Example Flow



Example

A common example is word counting. The mapper might count the occurrences of each word in a text document, and the reducer would then aggregate the counts from all the documents.

Suppose the input text is:

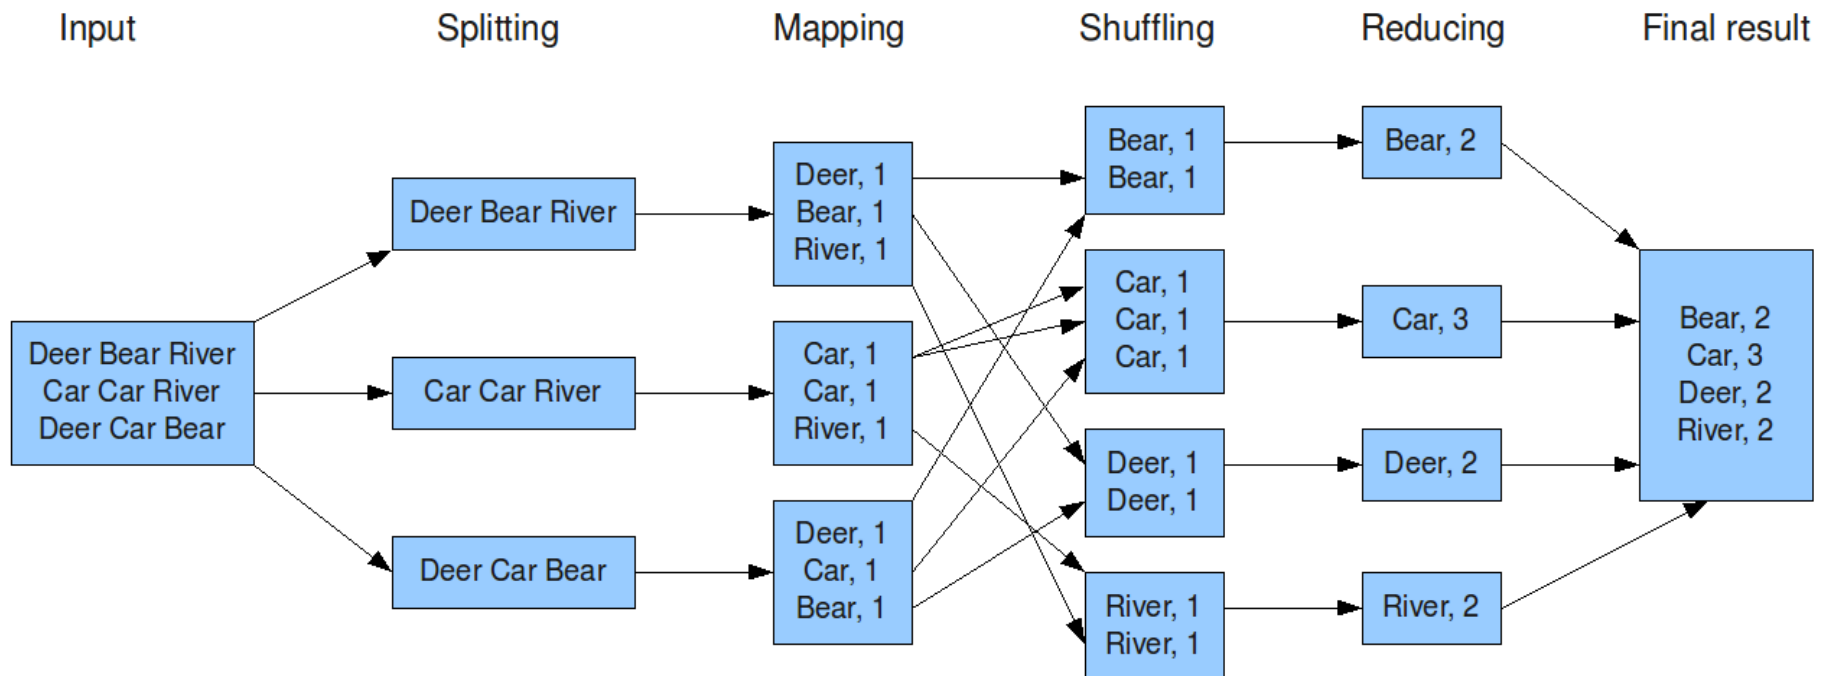
"Deer Bear River, Car Car River, Deer Car Bear".

It is split into chunks, each processed by mappers that output key-value pairs like ("Deer", 1). During the shuffle phase, all values for the same word are grouped together.

Finally, reducers aggregate these counts to produce the final word frequencies: Bear: 2, Car: 3, Deer: 2, River: 2.

Example

The overall MapReduce word count process



Map-Reduce

Advantages of using MapReduce

Scalability: MapReduce can handle very large datasets by distributing the processing load across multiple machines.

Parallelism: Map and Reduce functions allow for parallel processing, significantly reducing the overall processing time.

Fault Tolerance: MapReduce frameworks are designed to handle node failures, ensuring that processing can continue even if some machines go down.

Suitable for Unstructured Data: Can process **unstructured, semi-structured**, and structured data. Works with various data formats: text, JSON, XML, images, logs, etc.

Map-Reduce

Advantages of using MapReduce

Language Flexibility: Allows you to write Map and Reduce functions in multiple languages: Java, Python, C++, JavaScript (in MongoDB), etc.

Custom Processing Logic: User can define **custom business logic** using user-defined functions in Map and Reduce steps.
Suitable for complex workflows that traditional SQL can't express.

Data Locality Optimization: Moves **computation to where the data resides**, minimizing data transfer across the network.
In Hadoop, MapReduce jobs run on the same nodes that store the data in HDFS.

Ideal for: Large-scale **grouping** and **aggregation**, **Sorting** and **filtering**.

PPT Content Resources Reference Sample:

1. Book Reference

Meier, A., & Kaufmann, M. (2019). *SQL & NoSQL databases: Models, languages, consistency options, and architectures*. Springer. <https://link.springer.com/book/10.1007/978-3-658-24549-8>

2. Journal Article

Gupta, A., & Soni, S. (2018). A comparative study of NoSQL databases. *International Journal of Computer Sciences and Engineering*, 6(10), 123–130.

https://www.researchgate.net/publication/328095599_A_Comparative_study_of_NoSQL_Databases

3. Sources

Amazon Web Services. (n.d.). *Types of NoSQL databases*. Retrieved June 3, 2025, from <https://docs.aws.amazon.com/whitepapers/latest/choosing-an-aws-nosql-database/types-of-nosql-databases.html>

MongoDB. (n.d.). *What is NoSQL?* Retrieved June 3, 2025, from <https://www.mongodb.com/nosql-explained>

Google Cloud. (n.d.). *What is NoSQL?* Retrieved June 3, 2025, from <https://cloud.google.com/discover/what-is-nosql>

GeeksforGeeks. (n.d.). *Types of NoSQL databases*. Retrieved June 3, 2025, from <https://www.geeksforgeeks.org/types-of-nosql-databases/>

Parul[®]
University

NAAC
GRADE **A++**



<https://paruluniversity.ac.in/>



Nosql Data Management





CHAPTER-2

NoSQL Data Management



Introduction to NoSQL

What is NoSQL?

- NoSQL is non-relational database management system.
- A NoSQL database provides a mechanism for storage and retrieval of data that is modeled in means other than the tabular relations used in relational databases.
- It is a type of Non-Relational Database system.
- It follows CAP theorem rather than ACID properties of RDBMS.

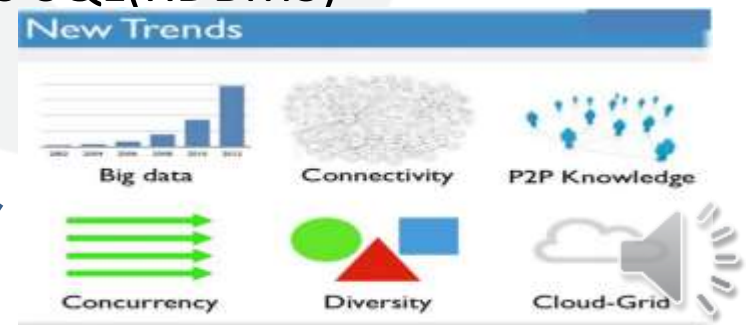




Why and when NoSql should be used ?

- The Traditional database is more predictable and preferable for Structured data.
- But, NoSQL is preferable for handling unstructured type of data . Unstructured data is type of data which has no proper format this type of data is generated by machine and humans in the form of satellite image, social media, mobile data, web-content etc
- NoSQL has high performance with high availability, and offers rich query language and easy scalability compared to SQL(RDBMS)

Why NoSQL is needed today most ???



Need of NoSQL

- Explosion of Social media sites (FaceBook, Twitter etc) resulting in generation of large data of different formats and size.
- Rise of Cloud based technologies like Amazon Web Services (AWS), Google cloud and Microsoft Azure
- Expansion of Open source communities.



Characteristics of NoSQL

- Does not follow ACID properties
- Less complexity as of SQL queries
- Schema less
- Fast development
- Large data volumes
- Easy and frequent DB changes



When to use NoSQL

- When Traditional RDBMS model provides restriction
- When ACID properties is not supported
- Temporary data (Like Wishlist, shopping cart, session data etc)
- When joins are expensive in RDBMS (product cost, hardware, maintenance)

When not to use NoSQL

- Financial Data, Data requiring strict ACID properties, Business critical data






CAP Theorem

- **Consistency**-This means that the data in the database remains consistent after the execution of an operation. For example after an update operation all clients see the same data.
- **Availability**-This means that the system is always on (service guarantee availability), no downtime.
- **Partition Tolerance**-This means that the system continues to function even the communication among the servers is unreliable, i.e. the servers may be partitioned into multiple groups that cannot communicate with one another. Here, if part of the database is unavailable, other parts are always unaffected.



CAP Theorem

- **CA** -Single site cluster, therefore all nodes are always in contact. When a partition occurs, the system blocks.
- **CP**-Some data may not be accessible, but the rest is still consistent/accurate.
- **AP**-System is still available under partitioning, but some of the data returned may be inaccurate.



The BASE

- The CAP theorem states that a distributed computer system cannot guarantee all of the following three properties at the same time:
 - Consistency
 - Availability
 - Partition tolerance
- A BASE system gives up on consistency.



The BASE

- **Basically Available** indicates that the system *does guarantee availability, in terms of the CAP theorem. i.e. DB is available all the time.*
- **Soft state** indicates that the state of the system may change over time, even without input. This is because of the eventual consistency model.
- **Eventual consistency** indicates that the system will become consistent over time, given that the system doesn't receive input during that time.





Eventual Consistency

- The term "eventual consistency" means to have copies of data on multiple machines to get high availability and scalability. Thus, changes made to any data item on one machine has to be propagated to other replicas.
- Data replication may not be instantaneous as some copies will be updated immediately while others in due course of time. These copies may be mutually, but in due course of time, they become consistent. Hence, the name eventual consistency.





Types of NoSQL

- There are 4 main types of NoSQL databases:

-Key value
-Column Family
-Document
-Graph





Key Value

- Data is stored in Key/value pairs. It is designed in such a way that it can handle lots of data and heavy load.
- Key-value pair storage databases store data as a hash table where each key is unique and the value can be JSON, BLOB(Binary Large Object), String etc can be accessed by string called keys
- For EX: A key value pair may contain key like “University” associated with value like “Parul”.

Key	Value
Name	Joe Bloggs
Age	42
Occupation	Stunt Double
Height	175cm
Weight	77kg





Key Value

- It is one of the most basic types of NoSQL databases
- This kind of database is used as a collection, dictionaries, associative arrays etc. key value stores help the developer to store schemaless data
- They work best for Shopping cart contents

Basic Operations are

Insert(Key,value), Fetch(key), Update(key), delete(Key).



Column- based

- Column oriented databases work on columns and are based on Big table paper by Google.
- The column is the lowest/ smallest instance of data.
- It is a tuple that contains a name, value and a timestamp
- Every column is treated separately.
- Values of single column databases are stored contiguously.

ColumnFamily			
Row Key	Column Name		
	Key	Key	Key
	Value	Value	Value
	Column Name		
	Key	Key	Key
	Value	Value	Value



Column based

- They deliver high performance on aggregation queries like SUM, COUNT, AVG, MIN etc. as the data is readily available in a column.
- Column-based NoSQL databases are widely used to manage data warehouses, business intelligence, CRM, Library card catalogs
- HBase, Cassandra, HBase, Hypertable are examples of column based database

Some statistics about Facebook search (Using Cassandra).

MySQL > 50gb data

Writes average – 300 ms

Reads average _ 350 ms

Rewritten in Cassandra < 50gb data

Writes average: 0.12 ms

reads average: 15ms





Document based

- Document-Oriented NoSQL DB stores and retrieves data as a key value pair but the value part is stored as a document.
- The document is stored in JSON or XML formats. The value is understood by the DB and can be queried.
- Pair each key with complex data structure known as data structure.
- Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

Col1	Col2	Col3	Col4
Data	Data	Data	Data
Data	Data	Data	Data
Data	Data	Data	Data

Document 1

```
{  
  "prop1": data,  
  "prop2": data,  
  "prop3": data,  
  "prop4": data  
}
```

Document 2

```
{  
  "prop1": data,  
  "prop2": data,  
  "prop3": data,  
  "prop4": data  
}
```

Document 3

```
{  
  "prop1": data,  
  "prop2": data,  
  "prop3": data,  
  "prop4": data  
}
```






Document based

- The document type is mostly used for CMS(content management system) systems, blogging platforms, real-time analytics & e-commerce applications.
- It should not use for complex transactions which require multiple operations or queries against varying aggregate structures.
- Amazon SimpleDB, CouchDB, MongoDB, Riak, Lotus Notes,are popular Document originated DBMS systems.



Graph based

- A graph type database stores entities as well the relations amongst those entities.
- The entity is stored as a node with the relationship as edges. An edge gives a relationship between nodes. Every node and edge has a unique identifier.





Graph based

- Compared to a relational database where tables are loosely connected, a Graph database is a multi-relational in nature. Traversing relationship is fast as they are already captured into the DB, and there is no need to calculate them.
- Graph base database mostly used for social networks, logistics, spatial data.
- Neo4J, Infinite Graph, OrientDB, FlockDB are some popular graph-based databases.





Comparison of SQL and NoSQL

1. Relational Database

Management System

2. SQL is also known as
Structured query language.

3. They have predefined
Schema

1. Non relational management
system

2. NoSQL is also known as Not
only Structured query
language

3. They have dynamic schema
for unstructured data



Comparison of SQL and NoSQL

4. Vertically scalable

5. RDBMS are Table based

6. SQL are better for multi-row transactions.

7. Example: Oracle, MySQL

4. Horizontally scalable

5. NoSQL databases are document, key-value, graph or wide-column based

6. NoSQL is better for unstructured data like documents or JSON

7. Example: MongoDB, Cassandra, CouchDB etc.





What is NewSQL

- The term NewSQL is not exactly as wide as NoSQL.
- NewSQL systems all begin with the relational data model and the SQL query language and they all attempt to cover portion of similar types of scalability, flexibility or lack-of-focus that has driven the NoSQL development.





Why NewSQL ?

- The early counter-measures to the above approaches were a powerful single node machine that is able to handle all the transaction, a custom built middleware system to distribute queries over traditional DBMS nodes.
- But both of them are prohibitively expensive to be carried out.





Why NewSQL ?

- As a result, there was a need of an intermediate database system which combines the distributed architectures of NoSQL systems with multiple node concurrency and a whole new storage mechanism.
- Thus Newsql can be defined as a class of modern relational DBMSs that seek to provide the same scalable performance of NoSQL for OLTP workloads and simultaneously guaranteeing ACID compliance for transactions as in RDBMS.
- these systems want to achieve the scalability of NoSQL without having to discard the relational model with SQL and transaction support of the legacy DBMS.





Concepts of NewSQL

- Main Memory storage of OLTP(Online transaction process) databases enables in-memory computations of databases.
- Scaling out by splitting a database into disjoint subsets called either partitions or shards, leading to the execution of a query into multiple partitions and then combine the result into a single result.
- NewSQL systems preserve the ACID properties of databases.
- Enhanced concurrency control system benefits upon traditional ones.



Concepts of NewSQL

- Presence of secondary index allowing NewSQL to support faster query processing times.
- High availability and Strong data durability is made possible with the use of replication mechanisms.
- NewSQL systems can be configured to provide synchronous updates of data over the WAN.
- Minimizes Downtime, provides fault tolerance with its crash recovery mechanism.





Comparison of NoSQL and NewSQL

- | | |
|--|---|
| 1. It doesn't follow a relational model, it was designed to be entirely different from that. | 1. Since the relational model is equally essential for real-time analytics. |
| 2. Supports CAP Theorem | 2. Supports ACID property |
| 3. Supports old SQL | 3. Supports old SQL and even enhanced functionality of old SQL |





Comparison of NoSQL and NewSQL

- 4. Supports OLTP database but it is not best suited.
- 5. Vertical Scaling
- 6. Better than SQL for processing complex queries
- 7. Supports distributed system

- 4. Supports OLTP database and highly efficient.
- 5. Vertical and horizontal scaling
- 6. Highly efficient to process complex queries and smaller queries
- 7. Supports distributed system





Map reduce

- Map-reduce is a data processing paradigm for condensing large volumes of data into useful aggregated results.
- MongoDB uses mapReduce command for map-reduce operations.
- MapReduce is generally used for processing large data sets.





MapReduce Command

```
>db.collection.mapReduce(  
  function() {emit(key,value);}, //map function  
  function(key,values) {return reduceFunction}, { //reduce function  
    out: collection,  
    query: document,  
    sort: document,  
    limit: number  
  }  
)
```





MapReduce Command

The map-reduce function first queries the collection, then maps the result documents to emit key-value pairs, which is then reduced based on the keys that have multiple values.

In the above syntax –

map is a javascript function that maps a value with a key and emits a key-value pair

reduce is a javascript function that reduces or groups all the documents having the same key

out specifies the location of the map-reduce query result

query specifies the optional selection criteria for selecting documents

sort specifies the optional sort criteria

limit specifies the optional maximum number of documents to be returned



Using MapReduce

Consider the following document structure storing user posts. The document stores user_name of the user and the status of post.

```
{  
  "post_text": "tutorialspoint is an awesome website for tutorials",  
  "user_name": "mark",  
  "status": "active"  
}
```





Using MapReduce

Now, we will use a mapReduce function on our posts collection to select all the active posts, group them on the basis of user_name and then count the number of posts by each user using the following code –

```
>db.posts.mapReduce(  
  function() { emit(this.user_id,1); },  
  
  function(key, values) {return Array.sum(values)}, {  
    query:{status:"active"},  
    out:"post_total"  
  }  
)
```



Using MapReduce

The above mapReduce query outputs the following result –

```
{  
  "result" : "post_total",  
  "timeMillis" : 9,  
  "counts" : {  
    "input" : 4,  
    "emit" : 4,  
    "reduce" : 2,  
    "output" : 2  
  },  
  "ok" : 1,  
}
```





Using MapReduce

The result shows that a total of 4 documents matched the query (status:"active"), the map function emitted 4 documents with key-value pairs and finally the reduce function grouped mapped documents having the same keys into 2.





```
>db.posts.mapReduce(
  function() { emit(this.user_id,1); },
  function(key, values) {return Array.sum(values)}, {
    query:{status:"active"},
    out:"post_total"
  }
).find()
```



Using MapReduce

The above query gives the following result which indicates that both users tom and mark have two posts in active states –

```
{ "_id" : "tom", "value" : 2 }  
{ "_id" : "mark", "value" : 2 }
```

In a similar manner, MapReduce queries can be used to construct large complex aggregation queries. The use of custom Javascript functions make use of MapReduce which is very flexible and powerful.



Aggregations

- Aggregations operations process data records and return computed results.
- Aggregation operations group values from multiple documents together, and can perform a variety of operations on the grouped data to return a single result.
- In SQL `count(*)` and `with group by` is an equivalent of MongoDB aggregation.



Aggregations

- For the aggregation in MongoDB, you should use aggregate() method.
- Syntax
- Basic syntax of aggregate() method is as follows –

```
>db.COLLECTION_NAME.aggregate(AGGREGATE_OPERATION)
```





Aggregations

- In the collection you have the following data –

```
{
  _id: ObjectId(7df78ad8902c)
  title: 'MongoDB Overview',
  description: 'MongoDB is no sql database',
  by_user: 'tutorials point',
  url: 'http://www.tutorialspoint.com',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 100
},
{
  _id: ObjectId(7df78ad8902d)
  title: 'NoSQL Overview',
  description: 'No sql database is very fast',
  by_user: 'tutorials point',
  url: 'http://www.tutorialspoint.com',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 10
},
{
  _id: ObjectId(7df78ad8902e)
  title: 'Neo4j Overview',
  description: 'Neo4j is no sql database',
  by_user: 'Neo4j',
  url: 'http://www.neo4j.com',
  tags: ['neo4j', 'database', 'NoSQL'],
  likes: 750
},
```





Aggregations

- Now from the above collection, if you want to display a list stating how many tutorials are written by each user, then you will use the following aggregate() method –

```
> db.mycol.aggregate([{$group : {_id : "$by_user", num_tutorial : {$sum : 1}}}])
{ "_id" : "tutorials point", "num_tutorial" : 2 }
{ "_id" : "Neo4j", "num_tutorial" : 1 }
>
```





Aggregations

- Sql equivalent query for the above use case will be select by_user, count(*) from mycol group by by_user.
- In the above example, we have grouped documents by field by_user and on each occurrence of by user previous value of sum is incremented.





Aggregations

Expression	Description	Example
\$sum	Sums up the defined value from all documents in the collection.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", num_tutorial : {\$sum : "\$likes"}}}])</code>
\$avg	Calculates the average of all given values from all documents in the collection.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", num_tutorial : {\$avg : "\$likes"}}}])</code>
\$min	Gets the minimum of the corresponding values from all documents in the collection.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", num_tutorial : {\$min : "\$likes"}}}])</code>





Aggregations

Expression	Description	Example
\$max	Gets the maximum of the corresponding values from all documents in the collection.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", num_tutorial : {\$max : "\$likes"}}}])</code>
\$push	Inserts the value to an array in the resulting document.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", url : {\$push: "\$url"}}}])</code>
\$addToSet	Inserts the value to an array in the resulting document but does not create duplicates.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", url : {\$addToSet : "\$url"}}}])</code>





Aggregations

Expression	Description	Example
\$first	Gets the first document from the source documents according to the grouping. Typically this makes only sense together with some previously applied "\$sort"-stage.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", first_url : {\$first : "\$url"}}}])</code>
\$last	Gets the last document from the source documents according to the grouping. Typically this makes only sense together with some previously applied "\$sort"-stage.	<code>db.mycol.aggregate([{\$group : {_id : "\$by_user", last_url : {\$last : "\$url"}}}])</code>





Aggregations

- In UNIX command, shell pipeline means the possibility to execute an operation on some input and use the output as the input for the next command and so on.
- MongoDB also supports same concept in aggregation framework. There is a set of possible stages and each of those is taken as a set of documents as an input and produces a resulting set of documents (or the final resulting JSON document at the end of the pipeline). This can then in turn be used for the next stage and so on.





Aggregations

- The MongoDB aggregation pipeline consists of stages. Each stage transforms the documents as they pass through the pipeline.
- Pipeline stages do not need to produce one output document for every input document; e.g., some stages may generate new documents or filter out documents.





Aggregations

- Following are the possible stages in aggregation framework –
- `$project` – Used to select some specific fields from a collection.
- `$match` – This is a filtering operation and thus this can reduce the amount of documents that are given as input to the next stage.
- `$group` – This does the actual aggregation as discussed above.
- `$sort` – Sorts the documents.
- `$skip` – With this, it is possible to skip forward in the list of documents for a given amount of documents.





Partitioning/Sharding

- Breaking large datasets into smaller ones and distributing datasets and query loads on those datasets are requisites to high scalability.
- If a dataset becomes very large for a single node or when high throughput is required, a single node cannot suffice.
- We need to partition/shard such datasets into smaller chunks and then each partition can act as a database on its own.



Partitioning

- Thus, a large dataset can be spread across many smaller partitions/shards and each can independently execute queries or run some programs.
- This way large executions can be parallelized across nodes(Partitions/Shards)



Why Partitioning ?

- The purpose behind partitioning is to spread data so that execution can be distributed across nodes.
- Along with partitioning, if we can ensure that every partition node takes a fair share, then at least in theory, 5 nodes should be able to handle 5 times as much data and 5 times as much read and write throughput of a single partition.
- If sharding is unfair, then a single node might be taking all the load and other nodes might sit idle. This defeats the purpose of sharding/partitioning. A good partition strategy should avoid Hot spots.





Partition by key-range

- Partition by key-range divides partitions based on certain ranges.
- For example, dividing an Organization based on the first letter of their name.
- So A-C, D-F, G-K, L-P, Q-Z is one of the ways by which whole Organization data can be partitioned in 5 parts.
- Now we know the boundaries of such partition so we can directly query a particular partition if we know the name of an employee.



Partition by key-range

- Ranges are not necessarily evenly spaced.
- As in the above example, partition Q-Z has 10 letters of the alphabet but partition A-C has only three.
- The reason behind such division is to allocate the same amount of data to different ranges.
- Due to the fact that most people have the name starting from letters between A-C, as compared to Q-Z, so this strategy will result in near equal distribution of data across the partition.





Partition by key-range

- One of the biggest benefits of such partitioning is the range queries.
- Suppose I need to find all people whose name starts from letters between R-S, then I only need to send the query to partition R-S.





Partition by key hash

- Key range partition strategy is quite prone to hot spots hence many distributed datastore employs another partitioning strategy.
- The hash value of the data's key is used to find out the partition. A good hash function can distribute data uniformly across multiple partitions.
- Cassandra, MongoDB, and Voldemort are databases employing a key hash-based strategy.



Partition by key hash

- Hash functions exercised for the purpose of partitioning should be cryptographically strong.
- Java's `Object.hashCode()` is usually avoided for this purpose as such implementation can lead to a large number of hash collisions.
- Each partition is then assigned a range of key hashes (Rather than range of keys) and all keys which fall within the parameter of partitions range will be stored on that partition. Partition ranges can be chosen to be evenly spaced or can be chosen from a strategy like consistent hashing.





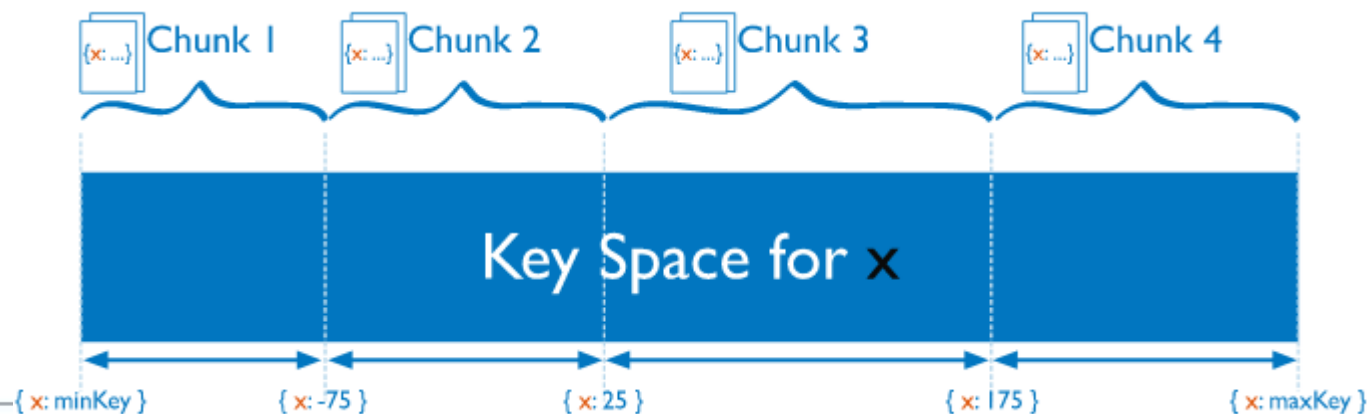
Partition by key hash

- Sadly with the hash of the key, we lose a nice property of key-range partitioning: efficient querying data with some ranges.
- As keys are now scattered to different partitions instead of being adjacent to a single partition.
- Hashing on key indeed reduces hot spots, however, it doesn't eliminate it completely. In case read and writes are for the same key, all requests still end up on the same partition.



Partitioning with Chunks

- MongoDB uses the shard key associated to the collection to partition the data into chunks.
- A chunk consists of a subset of sharded data.
- Each chunk has a inclusive lower and exclusive upper range based on the shard key.



Partitioning with Chunks

- MongoDB splits chunks when they grow beyond the configured chunk size. Both inserts and updates can trigger a chunk split.
- The smallest range a chunk can represent is a single unique shard key value. A chunk that only contains documents with a single shard key value cannot be split.



Partitioning with Chunks

- Initial Chunks
- The sharding operation creates the initial chunk(s) to cover the entire range of the shard key values. The number of chunks created depends on the configured chunk size.
- After the initial chunk creation, the balancer migrates these initial chunks across the shards as appropriate as well as manages the chunk distribution going forward.





Partitioning with Chunks

- Chunk Size
- The default chunk size in MongoDB is 64 megabytes. You can increase or reduce the chunk size. Consider the implications of changing the default chunk size:






Partitioning with Chunks

- Small chunks lead to a more even distribution of data at the expense of more frequent migrations. This creates expense at the query routing (mongos) layer.
- Large chunks lead to fewer migrations. This is more efficient both from the networking perspective and in terms of internal overhead at the query routing layer. But, these efficiencies come at the expense of a potentially uneven distribution of data.



Partitioning with Chunks





Partitioning with Chunks

- Chunk Migration
- MongoDB migrates chunks in a sharded cluster to distribute the chunks of a sharded collection evenly among shards. Migrations may be either:
- Manual. Only use manual migration in limited cases, such as to distribute data during bulk inserts.
- Automatic. The balancer process automatically migrates chunks when there is an uneven distribution of a sharded collection's chunks across the shards.






Partitioning with Chunks

- Balancing
- The balancer is a background process that manages chunk migrations.
- If the difference in number of chunks between the largest and smallest shard exceed the migration thresholds, the balancer begins migrating chunks across the cluster to ensure an even distribution of data.



Partitioning with Chunks



× ○ DIGITAL LEARNING CONTENT



Parul[®] University



www.paruluniversity.ac.in



NewSQL

Chapter-2: NoSQL Data Management

Mr. Janmejay Kumar Vishwakarma
Assistant Professor
Department of CSE

Content

1. Introduction to NoSQL.....	1
2. Why NoSQL?.....	3
3. CAP Theorem.....	4
4. BASE Model.....	8
5. Types of NoSQL.....	10
6. Advantages of NoSQL.....	20

NewSQL

- NewSQL is a class of modern relational database management.
- Support SQL and ACID transactions for online transaction processing (**OLTP**) read-write workloads.
 - Have **many small, fast transactions**
 - Access only a small part** of the data using indexes
 - Avoid full table scans** or big joins
 - Run **repetitive queries** with different input values

Key Features/ Advantages

SQL support (like traditional RDBMS).

ACID compliance (ensures data reliability).

High scalability (like NoSQL).

Optimized for **read-write intensive operations**

NewSQL

Why NewSQL?

Traditional RDBMS: Reliable but not scalable.

NoSQL: Scalable but compromises on consistency.

NewSQL offers the best of both worlds.

Popular NewSQL Databases:

- Google Spanner(Paypal, Spotify, Snapchat, X)
- CockroachDB(Netflix, Cisco)
- VoltDB (Huawei, Dream11)
- TiDB (Shopee, Bank of Beijing)

Use Cases

- E-commerce platforms
- Banking and finance systems
- Real-time analytics
- Large-scale web applications

Differences: SQL vs NoSQL vs NewSQL

Feature	SQL (Relational)	NoSQL (Non-relational)	NewSQL
Data Model	Relational tables (fixed schema)	Flexible: Key-value, document, column, graph (flexible)	Relational (adapted for distributed systems)
Scalability	Vertical (scale-up)	Horizontal (scale-out)	Horizontal + Vertical
ACID Support	Full ACID	Limited (BASE, eventual consistency)	Full ACID in distributed setup
SQL Support	Full standard SQL	No standard SQL	Full SQL + distributed extensions
Use Case	OLTP (small-scale, reliable)	Big Data, real-time, flexible workloads	Distributed OLTP with consistency & speed

Map-Reduce

- **MapReduce** is a programming model and an associated implementation for processing and generating big data sets with a parallel and distributed algorithm on a cluster.
- It is **highly scalable and reliable** and is **based on the principle of divide-and-conquer**, which provides built-in fault tolerance and redundancy.
- The MapReduce programming model was developed at Google in 2004 by Jeffrey Dean and Sanjay Ghemawat.

Map-Reduce Tasks

- It involves two key stages: **mapping** and **reducing**.

Mapping: function that processes input data and emits intermediate key-value pairs.

Reducing: function that aggregates or summarizes all values associated with the same key.

- MapReduce operations along with their individual stages:

Map task

- map
- combine (optional)
- partition

Reduce task

- shuffle and sort
- reduce

Map-Reduce Tasks

Map Tasks (Executed on Mapper Nodes)

1.Map: The main transformation stage.

- Input data is split into records and processed to generate intermediate **key-value pairs**.

Example: "apple orange apple" → ("apple", 1), ("orange", 1), ("apple", 1)

2.Combine (Optional):

- A **mini-reducer** that runs **locally** on the mapper node.
- Performs partial aggregation to reduce the amount of data transferred across the network.

Example: Combine ("apple", 1), ("apple", 1) → ("apple", 2) on the same node.

3.Partition:

- Decides **which reducer** each key-value pair should go to.
- Ensures all values for a given key are sent to the same reducer.
- Often uses a **hash function** on the key.

Map-Reduce Tasks

Reduce Tasks (Executed on Reducer Nodes)

1.Shuffle and Sort:

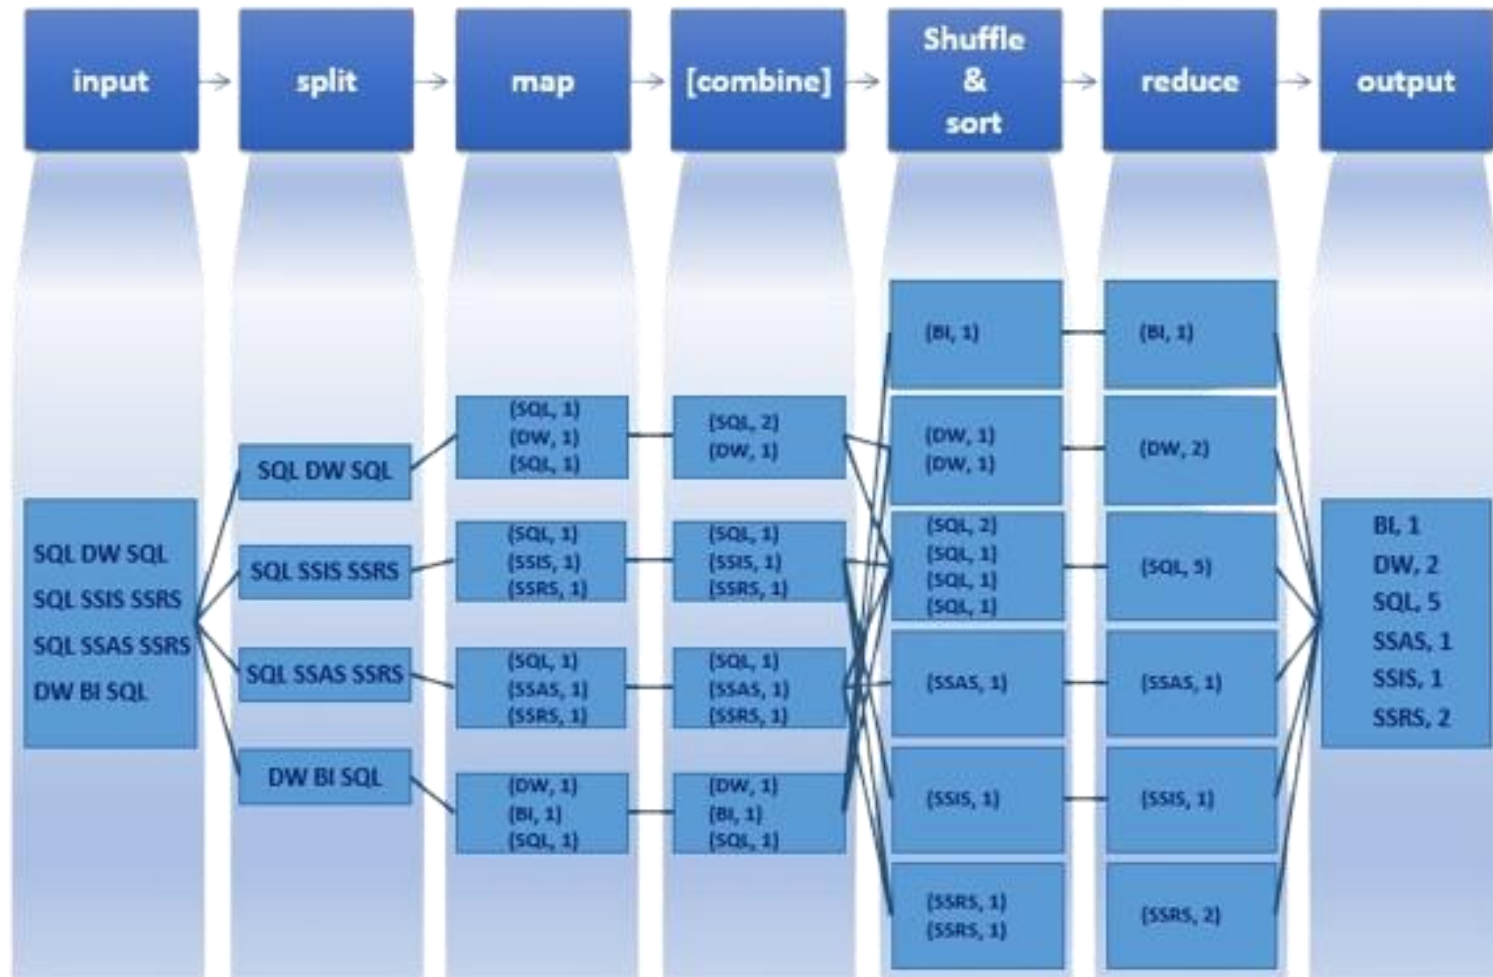
- Intermediate key-value pairs from all mappers are **shuffled** (transferred) across the network to appropriate reducers.
- The keys are then **sorted** before being passed to the reduce function.
- Ensures that a reducer gets all values for a single key grouped together.

2.Reduce:

- Applies user-defined logic on each group of values for a key.
- Produces the **final key-value output**.
- Example: ("apple", [1, 1, 2]) → ("apple", 4)

Example

MapReduce – Word Count Example Flow



Example

A common example is word counting. The mapper might count the occurrences of each word in a text document, and the reducer would then aggregate the counts from all the documents.

Suppose the input text is:


"Deer Bear River, Car Car River, Deer Car Bear".

It is split into chunks, each processed by mappers that output key-value pairs like ("Deer", 1). During the shuffle phase, all values for the same word are grouped together.

Finally, reducers aggregate these counts to produce the final word frequencies: Bear: 2, Car: 3, Deer: 2, River: 2.

Example

The overall MapReduce word count process



Map-Reduce

Advantages of using MapReduce

Scalability: MapReduce can handle very large datasets by distributing the processing load across multiple machines.

Parallelism: Map and Reduce functions allow for parallel processing, significantly reducing the overall processing time.

Fault Tolerance: MapReduce frameworks are designed to handle node failures, ensuring that processing can continue even if some machines go down.

Suitable for Unstructured Data: Can process **unstructured, semi-structured**, and structured data. Works with various data formats: text, JSON, XML, images, logs, etc.

Map-Reduce

Advantages of using MapReduce

Language Flexibility: Allows you to write Map and Reduce functions in multiple languages: Java, Python, C++, JavaScript (in MongoDB), etc.

Custom Processing Logic: User can define **custom business logic** using user-defined functions in Map and Reduce steps.
Suitable for complex workflows that traditional SQL can't express.

Data Locality Optimization: Moves **computation to where the data resides**, minimizing data transfer across the network.
In Hadoop, MapReduce jobs run on the same nodes that store the data in HDFS.

Ideal for: Large-scale **grouping** and **aggregation**, **Sorting** and **filtering**.

PPT Content Resources Reference Sample:

1. Book Reference

Meier, A., & Kaufmann, M. (2019). *SQL & NoSQL databases: Models, languages, consistency options, and architectures*. Springer. <https://link.springer.com/book/10.1007/978-3-658-24549-8>

2. Journal Article

Gupta, A., & Soni, S. (2018). A comparative study of NoSQL databases. *International Journal of Computer Sciences and Engineering*, 6(10), 123–130.

https://www.researchgate.net/publication/328095599_A_Comparative_study_of_NoSQL_Databases

3. Sources

Amazon Web Services. (n.d.). *Types of NoSQL databases*. Retrieved June 3, 2025, from <https://docs.aws.amazon.com/whitepapers/latest/choosing-an-aws-nosql-database/types-of-nosql-databases.html>

MongoDB. (n.d.). *What is NoSQL?* Retrieved June 3, 2025, from <https://www.mongodb.com/nosql-explained>

Google Cloud. (n.d.). *What is NoSQL?* Retrieved June 3, 2025, from <https://cloud.google.com/discover/what-is-nosql>

GeeksforGeeks. (n.d.). *Types of NoSQL databases*. Retrieved June 3, 2025, from <https://www.geeksforgeeks.org/types-of-nosql-databases/>

Parul[®]
University

NAAC
GRADE **A++**



<https://paruluniversity.ac.in/>

