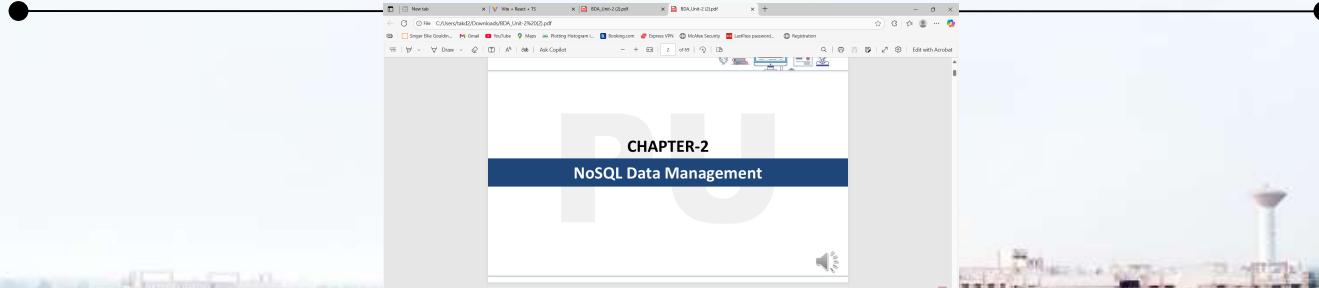




Nosql Data Management





A screenshot of a web browser window. The title bar says "New tab" and "File - Read - 15". The address bar shows "C:\Users\Acer\Downloads\BDA_Unit-2\BDA_Unit-2.pdf". The main content area displays a slide titled "CHAPTER-2" and "NoSQL Data Management". At the bottom right of the slide, there is a "DIGITAL LEARNING CONTENT" button with a play icon. The browser interface includes standard controls like back, forward, and search.

NoSQL Data Management

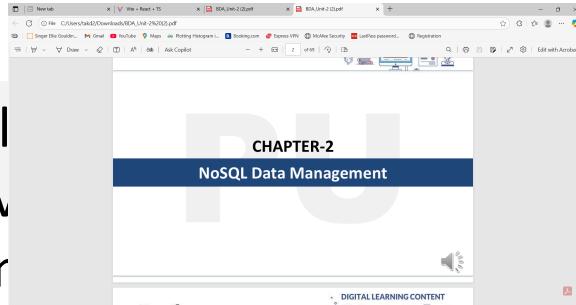




Introduction to NoSQL

What is NoSQL?

- NoSQL is non-relational
- A NoSQL database provides data that is modeled in ways that are different from relational databases.
- It is a type of Non-Relational Database system.
- It follows CAP theorem rather than ACID properties of RDBMS.



system.

storage and retrieval of data without the need for the complex, hierarchical relations used in relational databases.

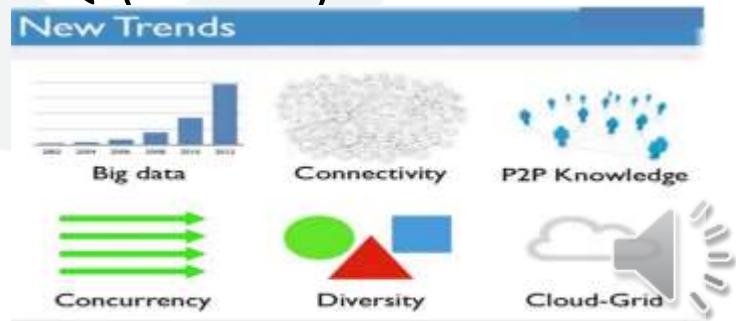




Why and when NoSql should be used ?

- The Traditional database is more predictable and preferable for Structured data.
- But, NoSQL is preferred for Unstructured data as most of the data is generated from unstructured sources like image, social media, mobile data, web-content etc
- NoSQL has high performance with high availability, and offers rich query language and easy scalability compared to SQL(RDBMS)

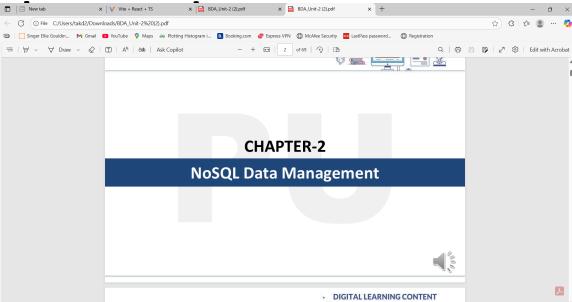
Why NoSQL is needed today most ???





Need of NoSQL

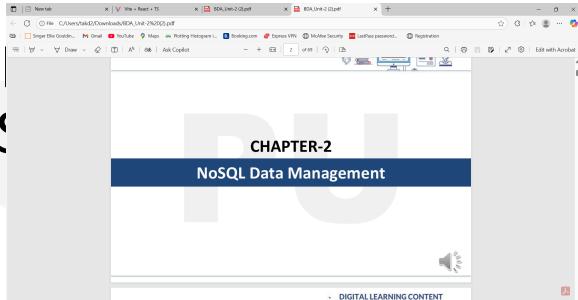
- Explosion of Social Media (FaceBook, Twitter etc.) leading to generation of large data in various formats and size.
- Rise of Cloud based technologies like Amazon Web Services (AWS), Google cloud and Microsoft Azure
- Expansion of Open source communities.





Characteristics of NoSQL

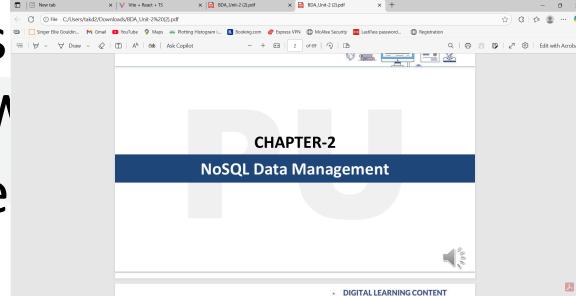
- Does not follow ACID
- Less complexity as of SQL
- Schema less
- Fast development
- Large data volumes
- Easy and frequent DB changes





When to use NoSQL

- When Traditional RDBMS model provides restriction
- When ACID properties
- Temporary data(Like V session data etc)
- When joins are expensive, product cost, hardware, maintenance)



When not to use NoSQL

- Financial Data, Data requiring strict ACID properties, Business critical data





CAP Theorem

- **Consistency**-This means that the data in the database remains consistent after the execution of an update operation. All clients see the same data.
- **Availability**-This means high availability, no downtime.
- **Partition Tolerance**-This means that the system continues to function even if the communication among the servers is unreliable, i.e. the servers may be partitioned into multiple groups that cannot communicate with one another. Here, if part of the database is unavailable, other parts are always unaffected.

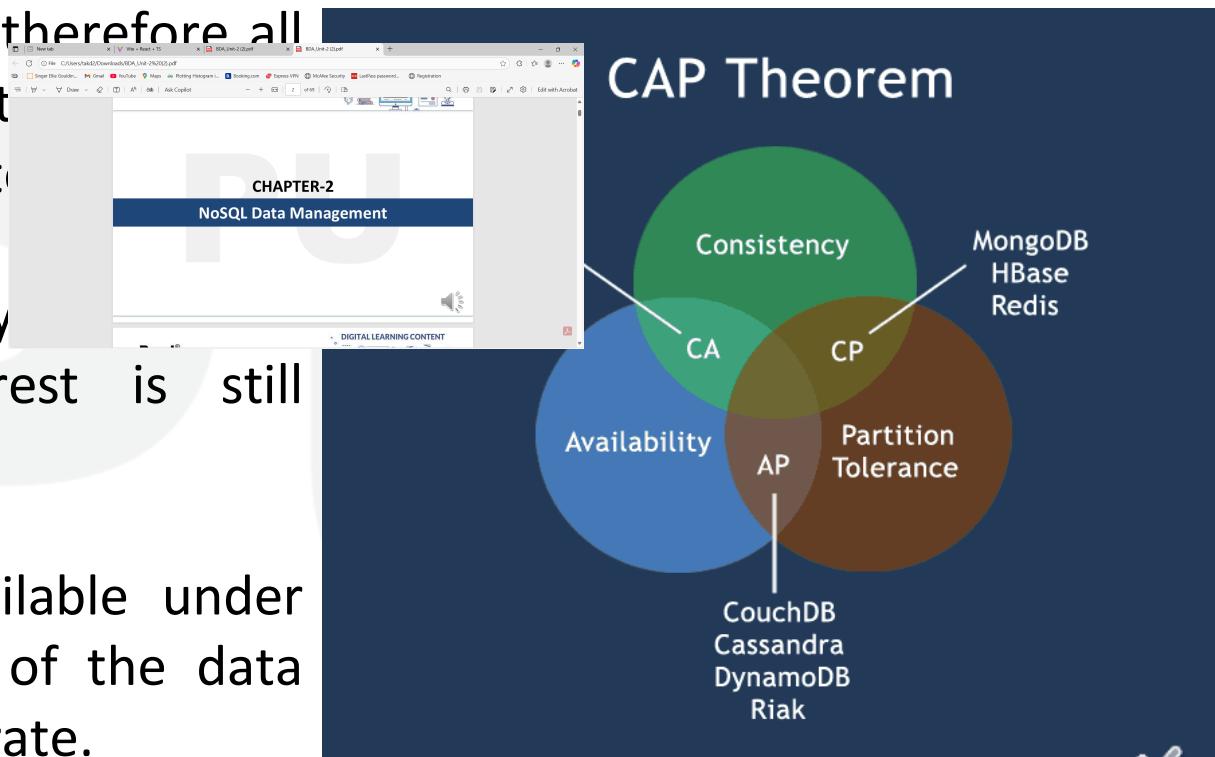


CAP Theorem

- **CA** -Single site cluster, therefore all nodes are always in contact. If a partition occurs, the system becomes unavailable.

- **CP**-Some data may be accessible, but the rest is still consistent/accurate.

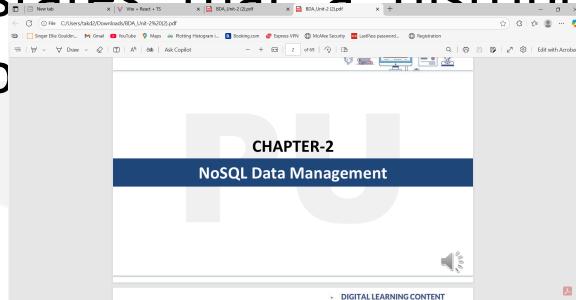
- **AP**-System is still available under partitioning, but some of the data returned may be inaccurate.





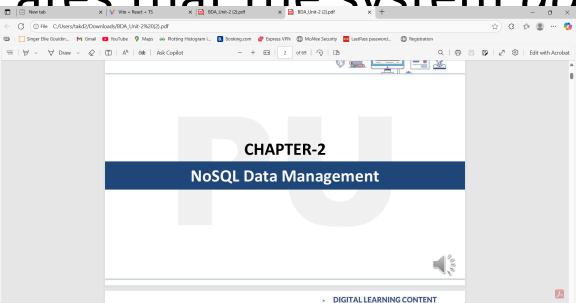
The BASE

- The CAP theorem states that a distributed computer system cannot guarantee all of the three properties at the same time:
 - Consistency
 - Availability
 - Partition tolerance
- A BASE system gives up on consistency.





The BASE

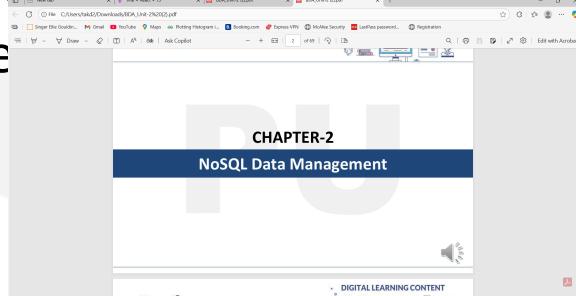
- **Basically Available** indicates that the system *does guarantee availability, in terms of the CAP theorem*.
A screenshot of a web browser window showing a digital learning content page. The title bar says 'New Tab' and 'File C:\Users\Shubham\Desktop\Digital Learning Content\Parul University\2nd Year\2nd Semester\UNIT-2\UNIT-2.pdf'. The main content area has a dark blue header with 'CHAPTER-2' and 'NoSQL Data Management' in white. Below the header is a large image of a book with the letters 'PU' on it. At the bottom of the page, there's a footer with the text 'DIGITAL LEARNING CONTENT'.
- **Soft state** indicates that even without input, the system's state
can change over time, indicating a *non-persistent* or *eventual consistency* model.
- **Eventual consistency** indicates that the system will become consistent over time, given that the system doesn't receive input during that time.





Eventual Consistency

- The term "eventual consistency" means to have copies of data on multiple machines to get scalability. Thus, changes made to any data item be propagated to other replicas.
- Data replication may not be instantaneous as some copies will be updated immediately while others in due course of time. These copies may be mutually, but in due course of time, they become consistent. Hence, the name eventual consistency.





Types of NoSQL

- There are 4 main types of NoSQL databases:

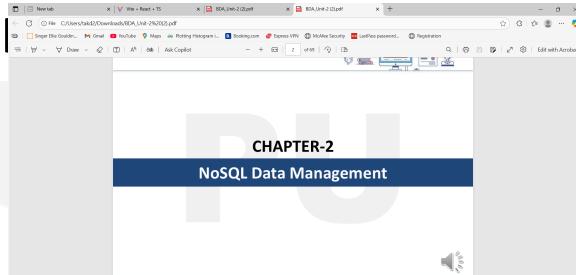
- Key value
- Column Family
- Document
- Graph





Key Value

- Data is stored in Key/value such a way that it can handle load.
- Key-value pair storage databases store data as a hash table where each key is unique and the value can be JSON, BLOB(Binary Large Object), String etc can be accessed by string called keys
- For EX: A key value pair may contain key like “University” associated with value like “Parul”.



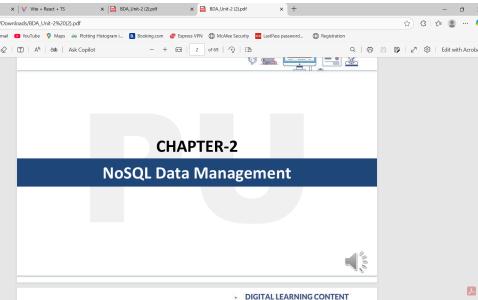
Key	Value
Name	Joe Bloggs
Age	42
Occupation	Stunt Double
Height	175cm
Weight	77kg





Key Value

- It is one of the most basic type of database.
- This kind of database is used to store key-value pairs, dictionaries, associative array.
- They help the developer to store schemaless data.
- They work best for Shopping cart contents.



Basic Operations are

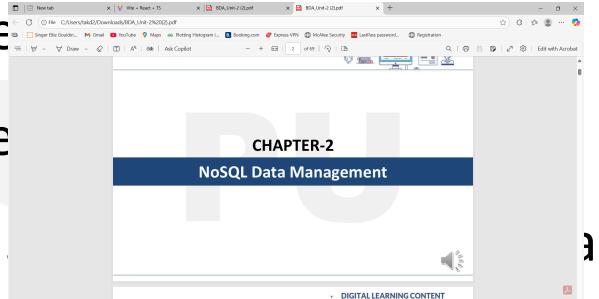
Insert(Key,value), Fetch(key), Update(key), delete(Key).





Column-based

- Column oriented database are based on Big table paper
 - The column is the lowest/
 - It is a tuple that contains a name, value and a timestamp
 - Every column is treated separately.
 - Values of single column databases are stored contiguously.



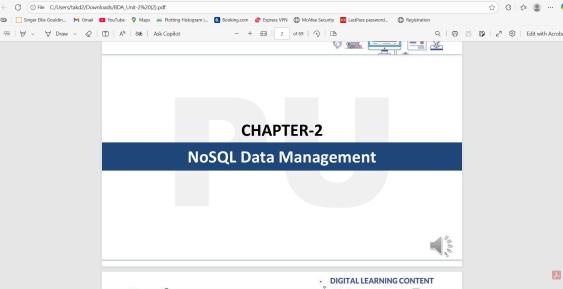
ColumnFamily			
Row Key	Column Name		
	Key	Key	Key
Value	Value	Value	Value
	Column Name		
Key	Key	Key	Key
	Value	Value	Value





Column based

- They deliver high performance on aggregation queries like SUM, COUNT, AVG, MIN etc. as the data is readily available.
- Column-based NoSQL databases are used for business intelligence, CRM, LinkedIn, manage data warehouses, etc.
- HBase, Cassandra, Hypertable are examples of column based database



manage data warehouses,

Some statistics about Facebook search(Using Cassandra).

MySQL > 50gb data

Writes average – 300 ms

Reads average _ 350 ms

Rewritten in Cassandra < 50gb data

Writes average: 0.12 ms

reads average: 15ms

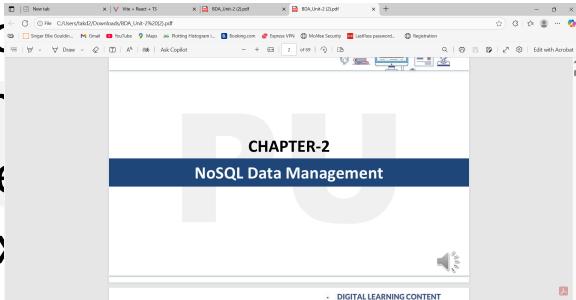




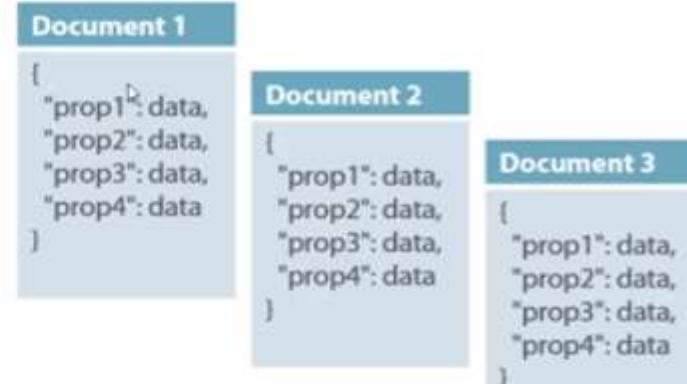
Document based

- Document-Oriented NoSQL DB stores and retrieves data as a key value pair but the value part is stored as a document.
- The document is stored in JSON format which is understood by the DB and can be queried.
- Pair each key with complex data structures.
- Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

Col1	Col2	Col3	Col4
Data	Data	Data	Data
Data	Data	Data	Data
Data	Data	Data	Data



The value is understood by the DB and can be queried. It is data structure.





Document based

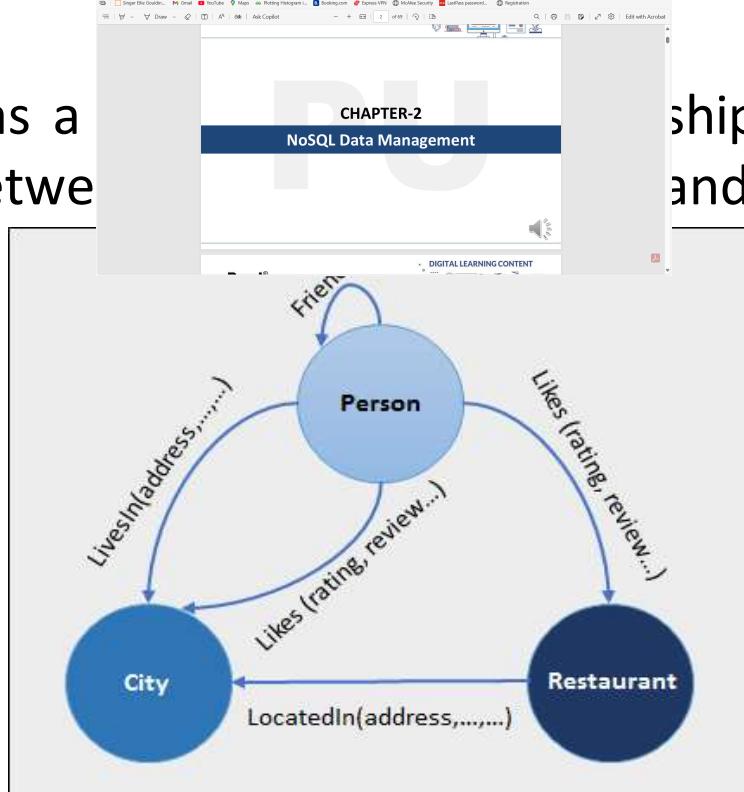
- The document type is mostly used for CMS/content management system) systems, blogging platforms, e-commerce applications.
- It should not use for operations or queries again which require multiple structures.
- Amazon SimpleDB, CouchDB, MongoDB, Riak, Lotus Notes, are popular Document originated DBMS systems.





Graph based

- A graph type database stores entities as well as the relations amongst those entities.
- The entity is stored as a node and relationship between them gives a relationship between entities.



ship as edges. An edge and edge has a unique identifier.





Graph based

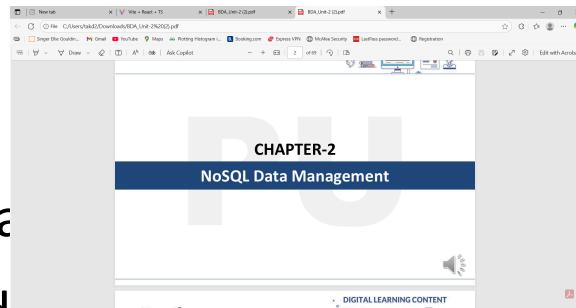
- Compared to a relational database where tables are loosely connected, a graph database is a multi-dimensional system where nodes are directly connected by edges. This makes it fast as they are already calculated and stored in memory, making them faster than relational databases when calculating complex queries. As they are already calculated and stored in memory, making them faster than relational databases when calculating complex queries.
- Graph base database models data as a graph structure, where nodes represent entities and edges represent relationships between entities. This makes it suitable for modeling complex relationships found in domains such as social networks, logistics, spatial data, and more.
- Neo4J, Infinite Graph, OrientDB, FlockDB are some popular graph-based databases.





Comparison of SQL and NoSQL

1. Relational Database Management System
2. SQL is also known as Structured query language
3. They have predefined Schema



1. Non relational management system
2. Also known as Not Structured query language
3. They have dynamic schema for unstructured data





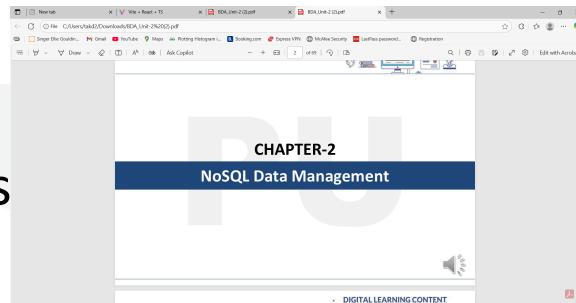
Comparison of SQL and NoSQL

4. Vertically scalable

5. RDBMS are Table bas

6. SQL are better for multi-
row transactions.

7. Example: Oracle, MySQL



4. Vertically scalable
5. RDBMS are document,
NoSQL databases are document,
graph or wide-column

based

6. NoSQL is better for unstructured
data like documents or JSON
7. Example: MongoDB, Cassandra,
CouchDB etc.



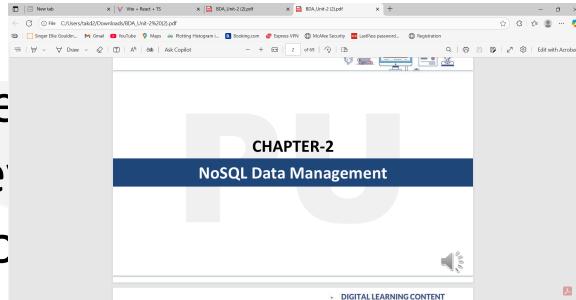


What is NewSQL

- The term NewSQL is not exactly as wide as NoSQL.

- NewSQL systems all benefit from the SQL query language and the ability to support scalability, flexibility and performance in development.

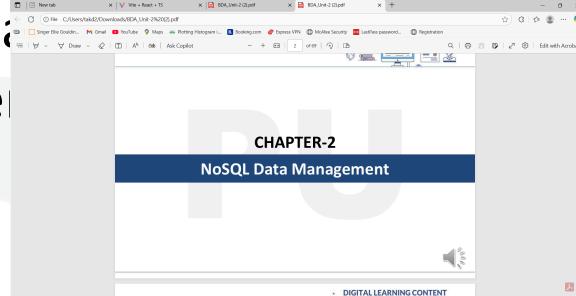
data model and the SQL portion of similar types of has driven the NoSQL





Why NewSQL ?

- The early counter-measures to the above approaches were a powerful single node machine that built middleware system over traditional DBMS nodes.
- But both of them are prohibitively expensive to be carried out.





Why NewSQL ?

- As a result, there was a need of an intermediate database system which combines the distributed nature of NoSQL systems with the ACID compliance mechanism of RDBMS.
- Thus Newsql can be considered as a class of DBMS that seek to provide the same scalable performance of NoSQL for OLTP workloads and simultaneously guaranteeing ACID compliance for transactions as in RDBMS.
- these systems want to achieve the scalability of NoSQL without having to discard the relational model with SQL and transaction support of the legacy DBMS.





Concepts of NewSQL

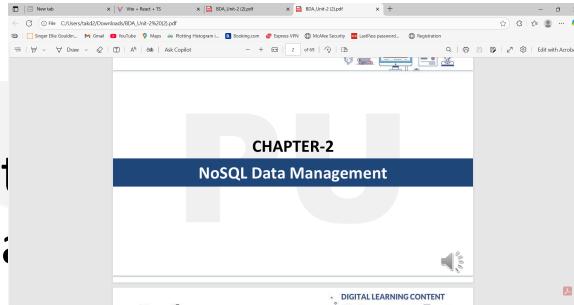
- Main Memory storage of OLTP(Online transition process) databases enables in-memory computation.
- Scaling out by splitting data into multiple partitions or shards, leading to parallel processing. This involves dividing data into subsets called either partitions or shards, executing different parts of a query into multiple partitions and then combining the results into a single result.
- NewSQL systems preserve the ACID properties of databases.
- Enhanced concurrency control system benefits upon traditional ones.





Concepts of NewSQL

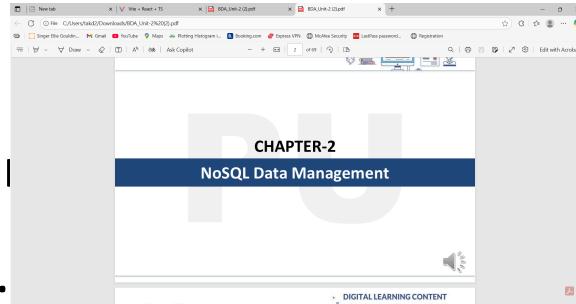
- Presence of secondary index allowing NewSQL to support faster query processing times.
- High availability and Strong consistency is made possible with the use of replication mechanism.
- NewSQL systems can be configured to provide synchronous updates of data over the WAN.
- Minimizes Downtime, provides fault tolerance with its crash recovery mechanism.





Comparison of NoSQL and NewSQL

1. It doesn't follow a relational model, it is designed to be entirely different from that.
2. Supports CAP Theorem
3. Supports old SQL



1. Since the relational model is essential for analytics.
2. Supports ACID property
3. Supports old SQL and even enhanced functionality of old SQL





Comparison of NoSQL and NewSQL

4. Supports OLTP database

but it is not best suited.

5. Vertical Scaling

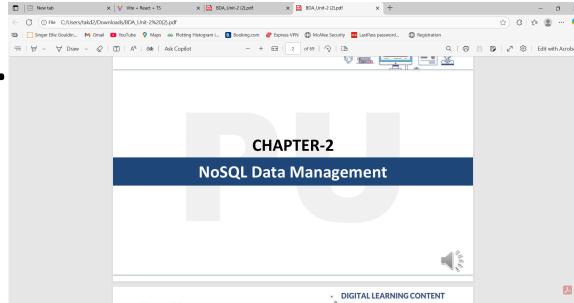
6. Better than SQL for
processing complex queries

7. Supports distributed system

4. Supports OLTP database

icient.

and horizontal



6. Highly efficient to process
complex queries and smaller
queries

7. Supports distributed system





Map reduce

- Map-reduce is a data processing paradigm for condensing large volumes of data into useful information.

- MongoDB uses mapReduce() function to perform map and reduce operations.

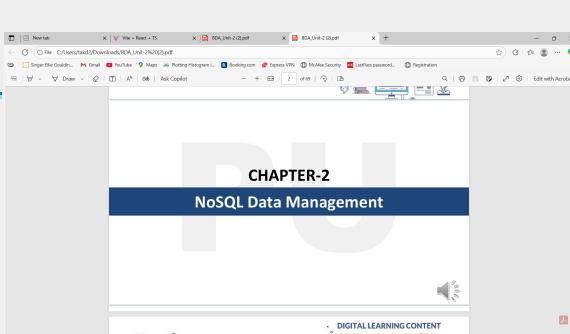
- MapReduce is generally used for processing large data sets.





MapReduce Command

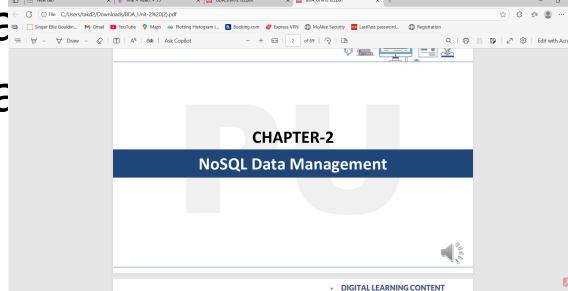
```
>db.collection.mapReduce(  
  function() {emit(key,v  
  function(key,values) {  
    out: collection,  
    query: document,  
    sort: document,  
    limit: number  
  }  
)  
  //reduce function
```





MapReduce Command

The map-reduce function first queries the collection, then maps the result documents to emit key-value pairs, which are then reduced based on the keys that have multiple values.



In the above syntax –

map is a javascript function that maps a value with a key and emits a key-value pair

reduce is a javascript function that reduces or groups all the documents having the same key

out specifies the location of the map-reduce query result

query specifies the optional selection criteria for selecting documents

sort specifies the optional sort criteria

limit specifies the optional maximum number of documents to be returned





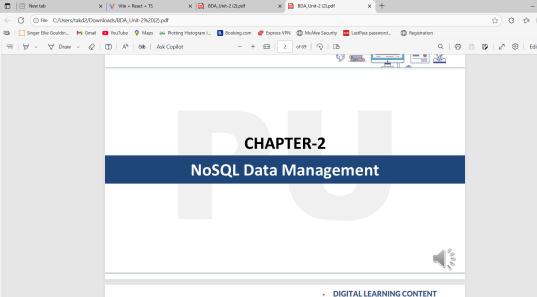
Using MapReduce

Consider the following document structure storing user posts. The document stores user_name and status of post.

{

```
"post_text": "tutorialsplatform",  
"user_name": "mark",  
"status": "active"
```

for tutorials",



}





Using MapReduce

Now, we will use a mapReduce function on our posts collection to select all the active posts basis of user_name and then count the number using the following code

```
>db.posts.mapReduce(  
  function() { emit(this.user_name, 1); },  
  
  function(key, values) { return Array.sum(values); }, {  
    query:{status:"active"},  
    out:"post_total"  
  }  
)
```

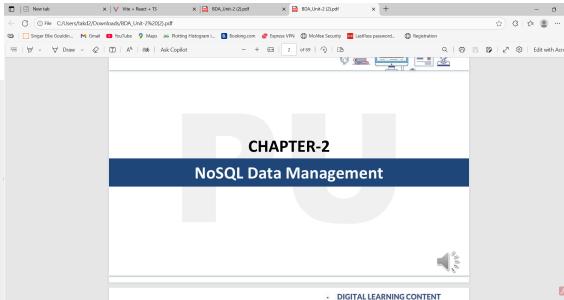




Using MapReduce

The above mapReduce query outputs the following result –

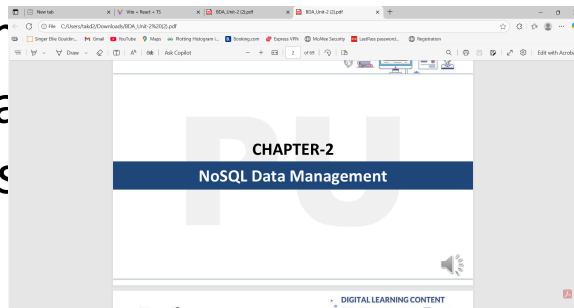
```
{  
    "result" : "post_total"  
    "timeMillis" : 9,  
    "counts" : {  
        "input" : 4,  
        "emit" : 4,  
        "reduce" : 2,  
        "output" : 2  
    },  
    "ok" : 1,  
}
```





Using MapReduce

The result shows that a total of 4 documents matched the query (status:"active"), the result contains 4 documents with key-value pairs and final output is shown below. The output is grouped mapped documents having the same key.

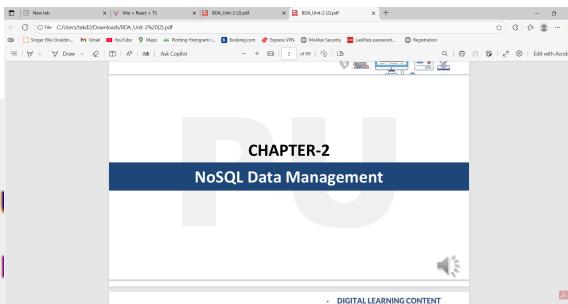




Using MapReduce

To see the result of this mapReduce query, use the find operator –

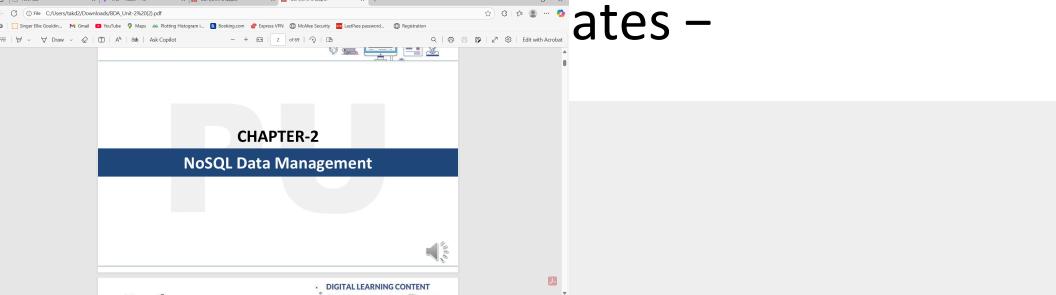
```
>db.posts.mapReduce(  
  function() { emit(this._id,  
    function(key, values) {  
      query:{status:"active"},  
      out:"post_total"  
    }  
  })  
.find()
```



Using MapReduce

The above query gives the following result which indicates that both users tom and mark have registered –

```
{ "_id" : "tom", "value" :  
{ "_id" : "mark", "value" :
```



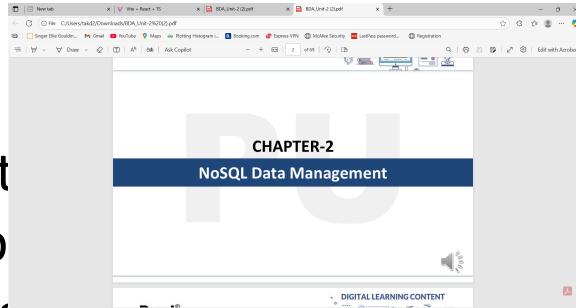
In a similar manner, MapReduce queries can be used to construct large complex aggregation queries. The use of custom Javascript functions make use of MapReduce which is very flexible and powerful.





Aggregations

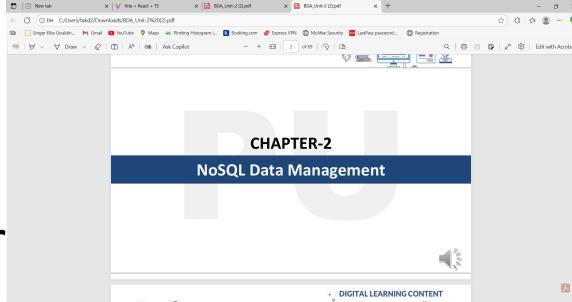
- Aggregations operations process data records and return computed results.
- Aggregation operations work together, and can perform multiple documents operations on the grouped data to return a single result.
- In SQL count(*) and with group by is an equivalent of MongoDB aggregation.





Aggregations

- For the aggregation in MongoDB, you should use aggregate() method.
- Syntax
- Basic syntax of aggr



OWS –

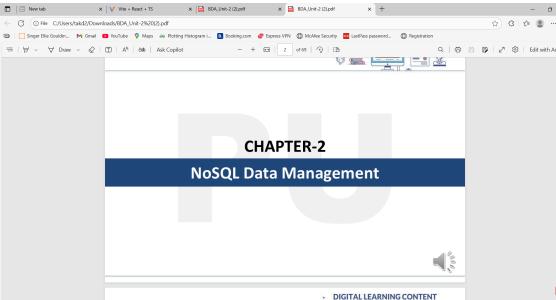
```
>db.COLLECTION_NAME.aggregate(AGGREGATE_OPERATION)
```





Aggregations

- In the collection you have the following data –



The screenshot shows a browser window with two tabs. The left tab displays a MongoDB document with the following data:

```
_id: ObjectId("7df78ad8902c")
title: 'MongoDB Overview',
description: 'MongoDB',
by_user: 'tutorials point',
url: 'http://www.tutorialspoint.com/mongodb/index.htm',
tags: ['mongodb', 'database'],
likes: 100
},
{
  _id: ObjectId("7df78ad8902d")
  title: 'NoSQL Overview',
  description: 'No sql database is very fast',
  by_user: 'tutorials point',
  url: 'http://www.tutorialspoint.com/nosql/index.htm',
  tags: ['mongodb', 'database', 'NoSQL'],
  likes: 10
},
{
  _id: ObjectId("7df78ad8902e")
  title: 'Neo4j Overview',
  description: 'Neo4j is no sql database',
  by_user: 'Neo4j',
  url: 'http://www.neo4j.com',
  tags: ['neo4j', 'database', 'NoSQL'],
  likes: 750
},
```

The right tab shows a digital learning content slide titled "CHAPTER-2 NoSQL Data Management". The slide has a dark blue header with the title and a white body. At the bottom right of the slide, there is a "DIGITAL LEARNING CONTENT" watermark.





Aggregations

- Now from the above collection, if you want to display a list stating how many tutorials each user, then you will use the following aggregation query:

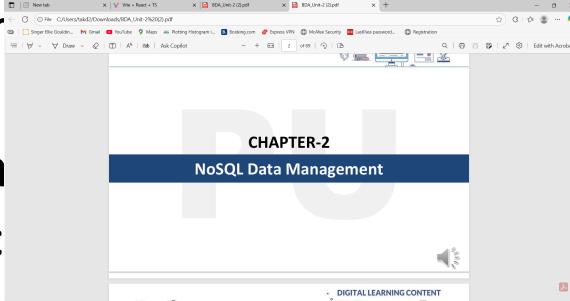
```
> db.mycol.aggregate([{$group : {_id : "$buy_user", numTutorial : {$sum : 1}}}] )
{ "_id" : "tutorials point", "numTutorial" : 2 }
{ "_id" : "Neo4j", "numTutorial" : 1 }
>
```





Aggregations

- Sql equivalent query for the above use case will be select by_user, count(*) from user.
- In the above example we are using group by_user and on each document by field sum is incremented.



I documents by field previous value of sum





Aggregations

Expression	Description	Example
\$sum	Sum from collection.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", numTutorial : {\$sum : "\$likes"}}}])</pre>
\$avg	Calculates the average of all given values from all documents in the collection.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", numTutorial : {\$avg : "\$likes"}}}])</pre>
\$min	Gets the minimum of the corresponding values from all documents in the collection.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", numTutorial : {\$min : "\$likes"}}}])</pre>





Aggregations

Expression	Description	Example
\$max	Gets the maximum value in a document.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", numTutorial : {\$max : "\$likes"}}}])</pre>
\$push	Inserts the value to an array in the resulting document.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", url : {\$push: "\$url"}}}])</pre>
\$addToSet	Inserts the value to an array in the resulting document but does not create duplicates.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", url : {\$addToSet : "\$url"}}}])</pre>





Aggregations

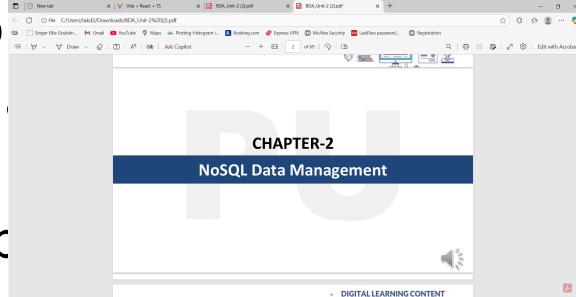
Expression	Description	Example
\$first	Gets the source document to the first stage.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", first_url : {\$first : "\$url"}}}])</pre>
\$last	Gets the last document from the source documents according to the grouping. Typically this makes only sense together with some previously applied “\$sort”-stage.	<pre>db.mycol.aggregate([{\$group : {_id : "\$by_user", last_url : {\$last : "\$url"}}}])</pre>





Aggregations

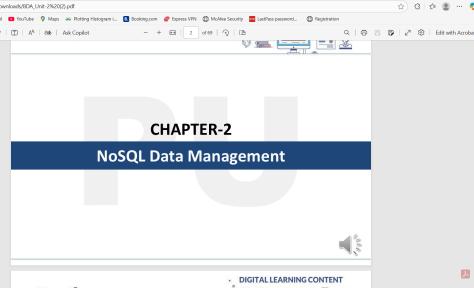
- In UNIX command, shell pipeline means the possibility to execute an operation on some output as the input for the next command
- MongoDB also supports aggregation framework. There is a set of possible stages and each of those is taken as a set of documents as an input and produces a resulting set of documents (or the final resulting JSON document at the end of the pipeline). This can then in turn be used for the next stage and so on.





Aggregations

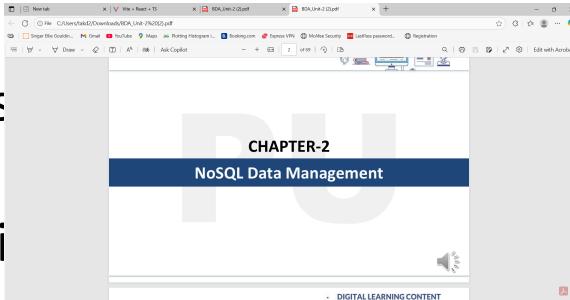
- The MongoDB aggregation pipeline consists of stages. Each stage transforms the document as it moves through the pipeline.
- Pipeline stages do not produce an output document for every input document. Stages may generate new documents or filter out documents.





Aggregations

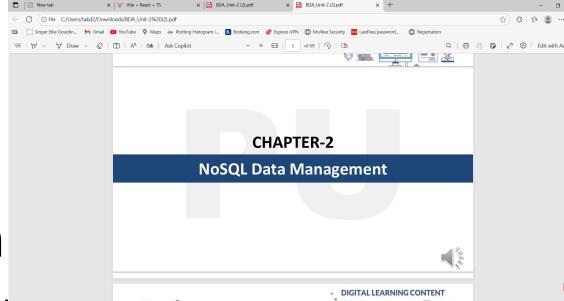
- Following are the possible stages in aggregation framework –
- \$project – Used to select fields from a collection.
- \$match – This is a filter stage. With this can reduce the amount of documents that are given as input to the next stage.
- \$group – This does the actual aggregation as discussed above.
- \$sort – Sorts the documents.
- \$skip – With this, it is possible to skip forward in the list of documents for a given amount of documents.





Partitioning/Sharding

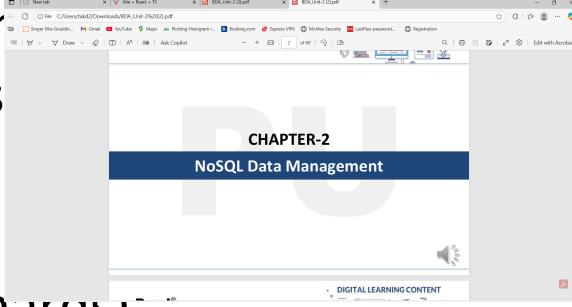
- Breaking large datasets into smaller ones and distributing datasets and query are requisites to high scalability.
- If a dataset becomes too large for a single node or when high throughput is required, a single node cannot suffice.
- We need to partition/shard such datasets into smaller chunks and then each partition can act as a database on its own.





Partitioning

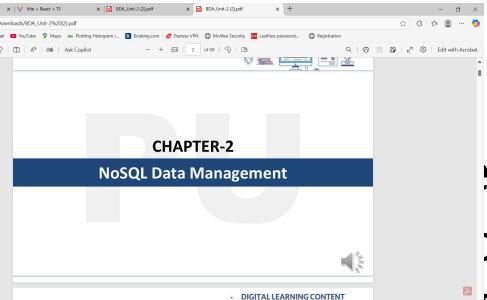
- Thus, a large dataset can be spread across many smaller partitions/shards and execute queries or run some programs
- This way large parallelized across nodes(Partitions/Shards)





Why Partitioning ?

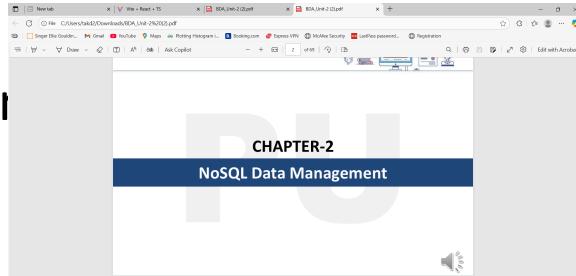
- The purpose behind partitioning is to spread data so that execution can be distributed among multiple nodes.
- Along with partitioning, it is also important to ensure that every partition has a fair share of data. For example, if there are 5 nodes in a cluster, 5 nodes should be able to handle 5 times as much data and 5 times as much read and write throughput of a single partition.
- If sharding is unfair, then a single node might be taking all the load and other nodes might sit idle. This defeats the purpose of sharding/partitioning. A good partition strategy should avoid Hot spots.





Partition by key-range

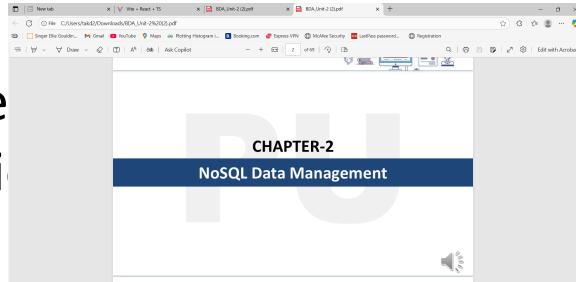
- Partition by key-range divides partitions based on certain ranges.
- For example, dividing data on the first letter of their name.
- So A-C, D-F, G-K, L-R, Q-Z is one of the ways by which whole Organization data can be partitioned in 5 parts.
- Now we know the boundaries of such partition so we can directly query a particular partition if we know the name of an employee.





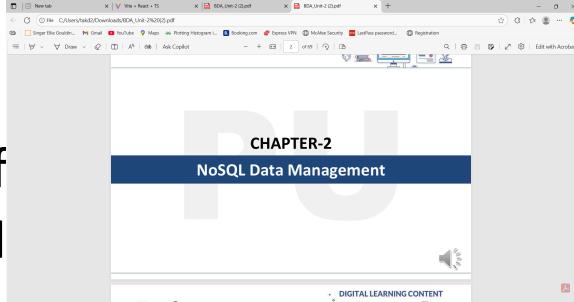
Partition by key-range

- Ranges are not necessarily evenly spaced.
- As in the above example, 'A' has 10 letters of the alphabet but partitioned into 3 ranges.
- The reason behind such division is to allocate the same amount of data to different ranges.
- Due to the fact that most people have the name starting from letters between A-C, as compared to Q-Z, so this strategy will result in near equal distribution of data across the partition.



Partition by key-range

- One of the biggest benefits of such partitioning is the range queries.
 - Suppose I need to find all documents whose key starts from letters R-S, then I can use a query to partition R-S.



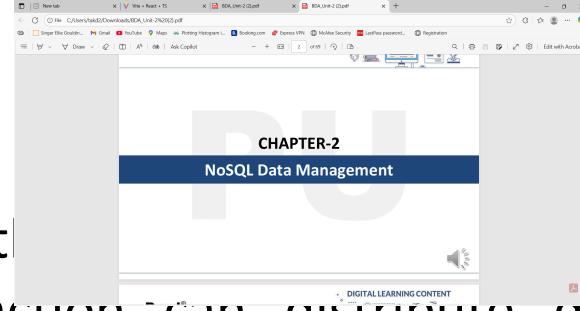
me starts from letters
uery to partition R-S.





Partition by key hash

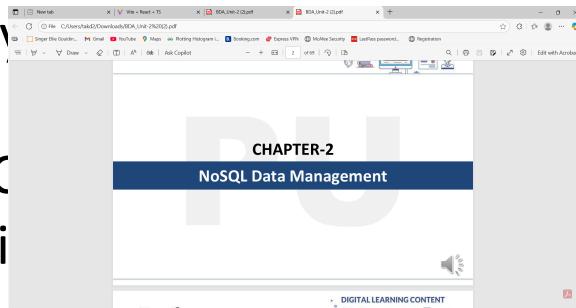
- Key range partition strategy is quite prone to hot spots hence many distributed systems have adopted another partitioning strategy.
- The hash value of the primary key is used to find out the partition. A good hash function can distribute data uniformly across multiple partitions.
- Cassandra, MongoDB, and Voldemort are databases employing a key hash-based strategy.





Partition by key hash

- Hash functions exercised for the purpose of partitioning should be cryptographically strong.
- Java's Object.hashCode() is not suitable for this purpose as it uses a large number of hash collisions.
- Each partition is then assigned a range of key hashes (Rather than range of keys) and all keys which fall within the parameter of partitions range will be stored on that partition. Partition ranges can be chosen to be evenly spaced or can be chosen from a strategy like consistent hashing.



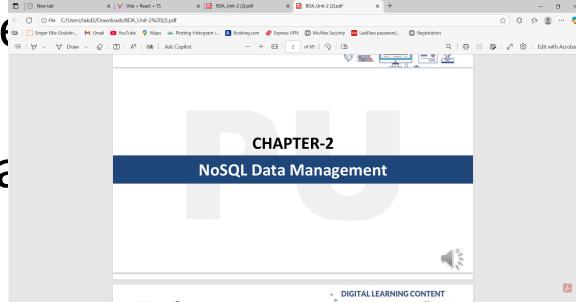
d for this purpose as
e number of hash





Partition by key hash

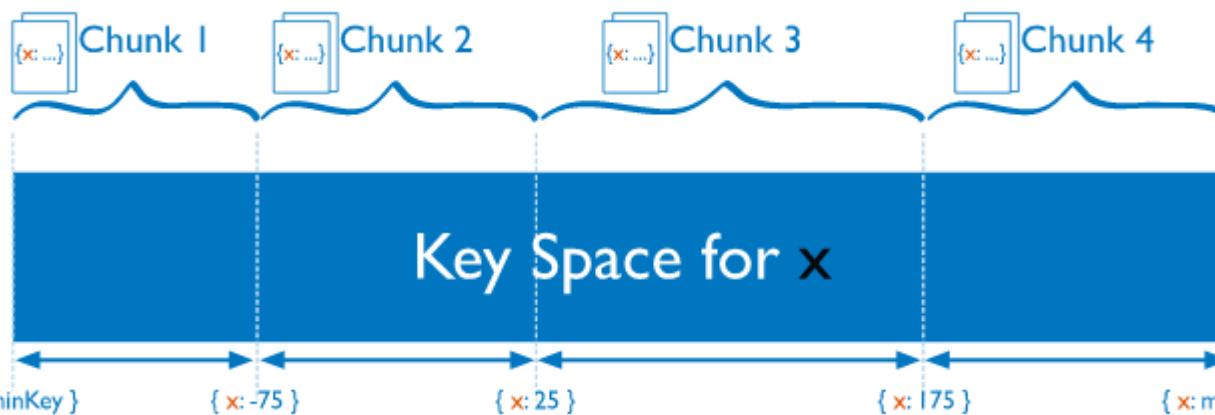
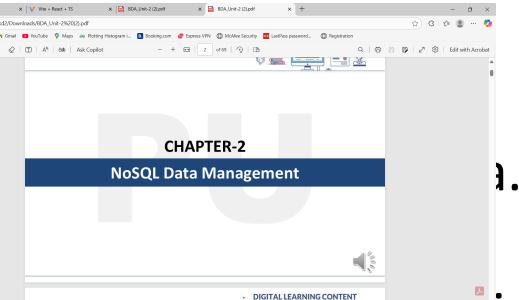
- Sadly with the hash of the key, we lose a nice property of key-range partitioning: each key ranges with some ranges.
- As keys are now scattered across partitions instead of being adjacent to a single partition.
- Hashing on key indeed reduces hot spots, however, it doesn't eliminate it completely. In case read and writes are for the same key, all requests still end up on the same partition.





Partitioning with Chunks

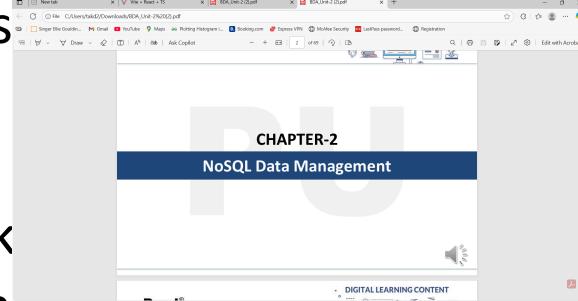
- MongoDB uses the shard key associated to the collection to partition the data into chunks.
- A chunk consists of documents.
- Each chunk has an inclusive lower and exclusive upper range based on the shard key.





Partitioning with Chunks

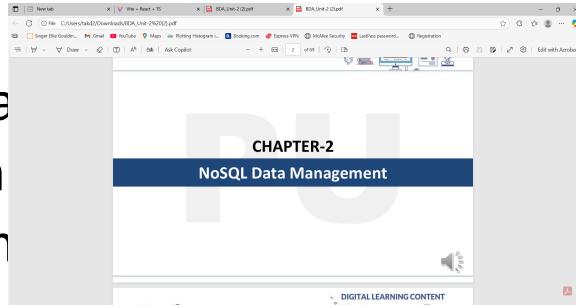
- MongoDB splits chunks when they grow beyond the configured chunk size. Both instances trigger a chunk split.
- The smallest range key value. A chunk with a single unique shard key value cannot be split.





Partitioning with Chunks

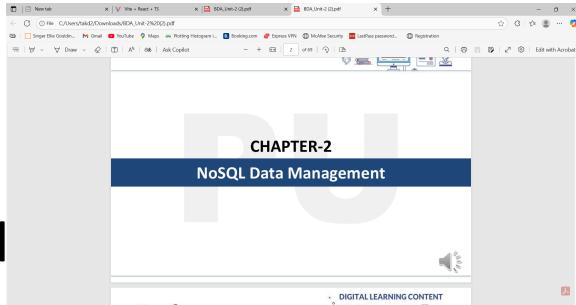
- Initial Chunks
- The sharding operation depends on the entire range of the data. The number of chunks created depends on the size of the data.
- After the initial chunk creation, the balancer migrates these initial chunks across the shards as appropriate as well as manages the chunk distribution going forward.





Partitioning with Chunks

- Chunk Size
- The default chunk size can be increased or reduced by changing the default



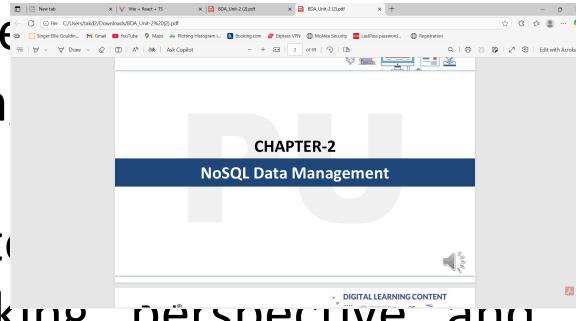
megabytes. You can consider the implications of





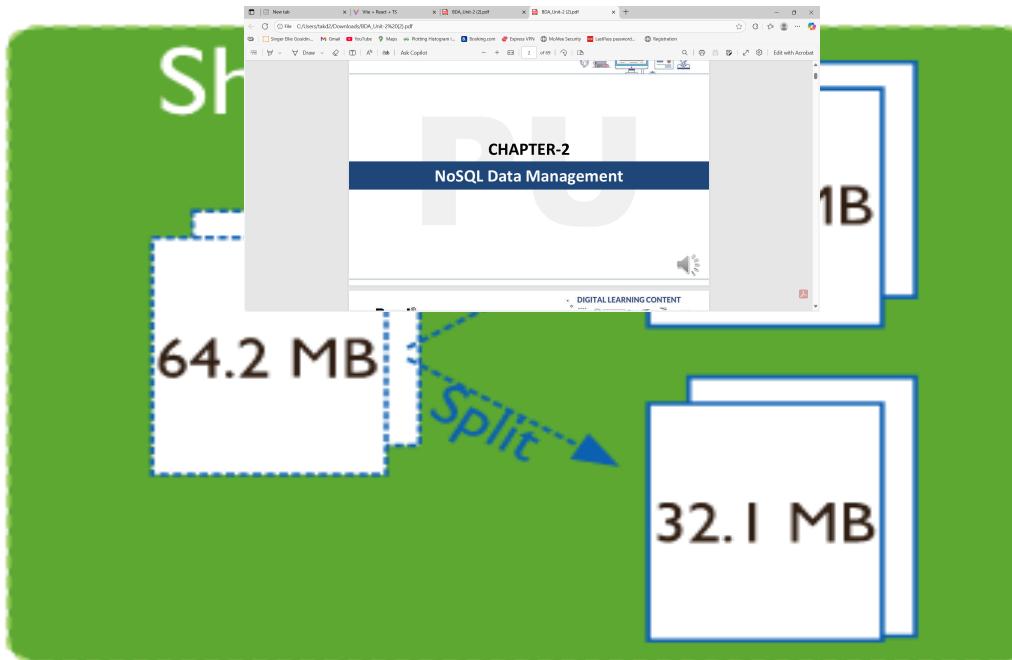
Partitioning with Chunks

- Small chunks lead to a more even distribution of data at the expense of more frequent query routing (monitors) creates expense at the expense of a potentially uneven distribution of data.
- Large chunks lead to less overhead from the networking perspective and in terms of internal overhead at the query routing layer. But, these efficiencies come at the expense of a potentially uneven distribution of data.





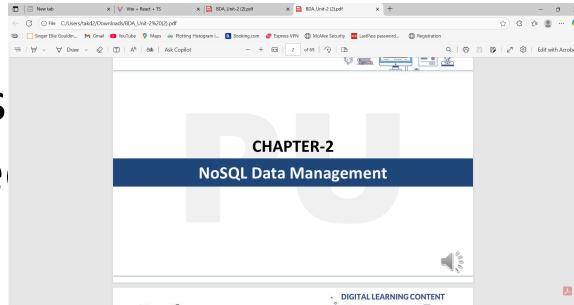
Partitioning with Chunks





Partitioning with Chunks

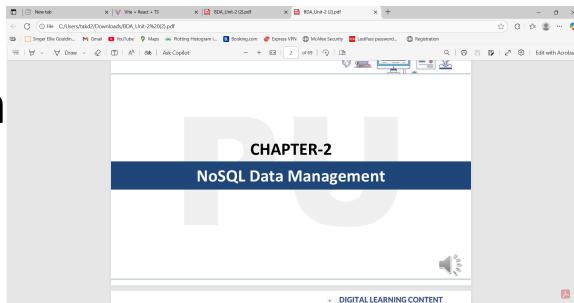
- Chunk Migration
- MongoDB migrates chunks of a shard. There may be either:
 - Manual. Only use manual migration in limited cases, such as to distribute data during bulk inserts.
 - Automatic. The balancer process automatically migrates chunks when there is an uneven distribution of a sharded collection's chunks across the shards.





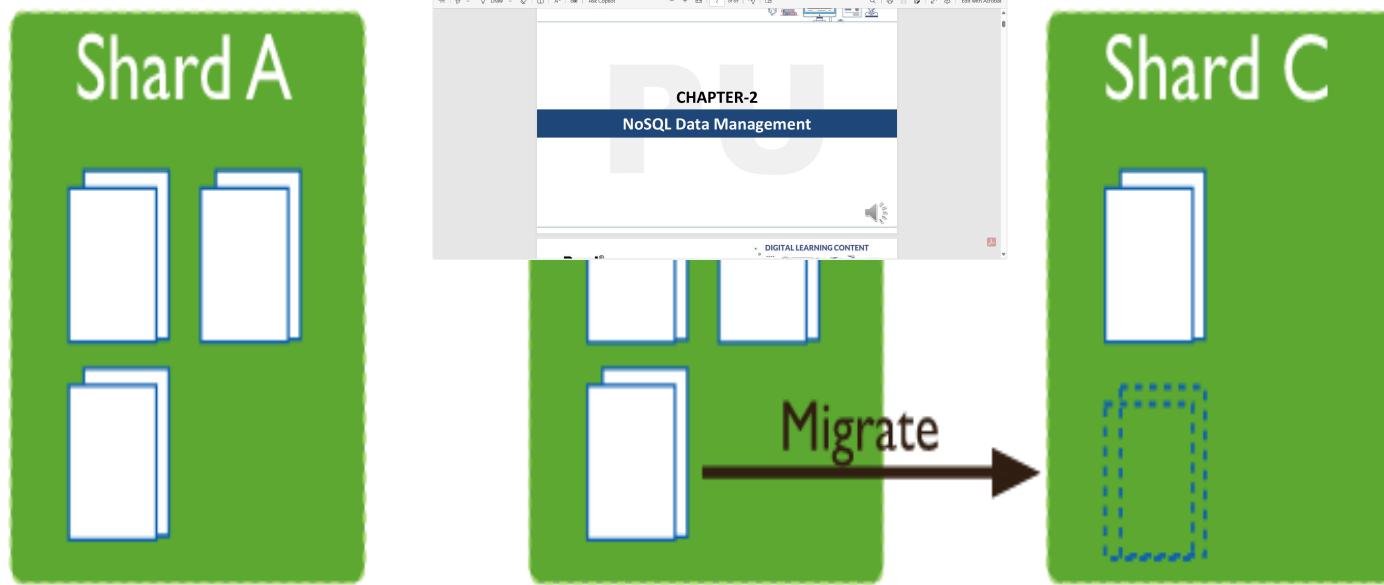
Partitioning with Chunks

- Balancing
- The balancer is a service that manages chunk migrations.
- If the difference in number of chunks between the largest and smallest shard exceed the migration thresholds, the balancer begins migrating chunks across the cluster to ensure an even distribution of data.

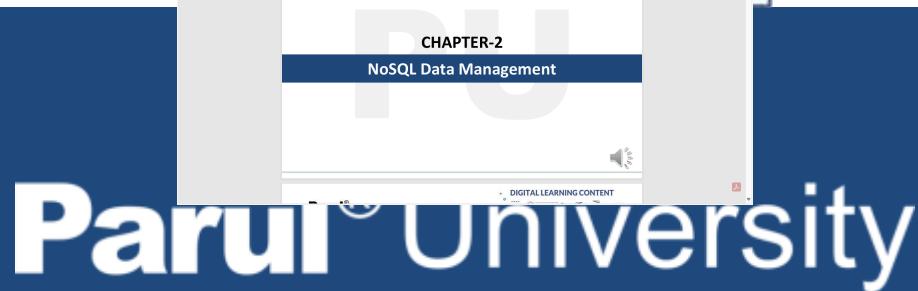




Partitioning with Chunks



DIGITAL LEARNING CONTENT



www.paruluniversity.ac.in

